

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 1

Not yet
answered

Marked out of
4.00

🚩 [Flag question](#)

Time left 1:29:49

Given a linear model with explicit bias term $g(\mathbf{x}; \mathbf{w}, b)$, we want to obtain an update rule for b similar to the obtained update rule for $\mathbf{w}$:

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \frac{\partial}{\partial \mathbf{w}} R_{emp}(\mathbf{y}, \mathbf{x}, \mathbf{w}, b).$$

$$\text{Note that } R_{emp}(\mathbf{y}, \mathbf{x}, \mathbf{w}, b) = \frac{1}{2} \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n + b))^2$$

Select one:

- ☐ a. $b^{new} = b^{old} - \eta \frac{1}{2} \sum_{n=1}^N ((\mathbf{w}^T \mathbf{x}^n + b) - y^n) \mathbf{x}^n$
- ☐ b. $b^{new} = b^{old} - \eta \frac{1}{2} \sum_{n=1}^N ((\mathbf{w}^T \mathbf{x}^n + b) - y^n)$
- ☐ c. $b^{new} = b^{old} - \eta \sum_{n=1}^N ((\mathbf{w}^T \mathbf{x}^n + b) - y^n) \mathbf{x}^n$
- ☐ d. $b^{new} = b^{old} - \eta \sum_{n=1}^N ((\mathbf{w}^T \mathbf{x}^n + b) - y^n)$

Next page

◀ [Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:40

Question 2

Not yet
answered

Marked out of
3.00

 [Flag question](#)

As described in the lecture, for single-layer neural networks with quadratic loss, we have the following objective function $R(b) = \frac{1}{2} \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n + b))^2$, where we made the bias parameter b explicit. Which bias parameter optimizes the objective function assuming fixed $\mathbf{w}$?

Select one:

- ☐ a. $b = \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n))$
- ☐ b. $b = \frac{1}{N} \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n))$
- ☐ c. $b = \frac{1}{2N} \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n))$
- ☐ d. $b = \frac{2}{N} \sum_{n=1}^N (y^n - (\mathbf{w}^T \mathbf{x}^n))$
- ☐ e. $b = \frac{1}{N} \sum_{n=1}^N (x^n - (\mathbf{w}^T \mathbf{y}^n))$

[Previous page](#)

[Next page](#)

[◀ Slides](#)



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:33

Question 3

Not yet answered

Marked out of 3.00

[Flag question](#)

What is the objective function of the perceptron?

Select one:

- ☐ a. $-\sum_{n=1, y^n=1}^N y^n \mathbf{w}^T \mathbf{x}^n$
- ☐ b. $-\sum_{n=1, y^n=-1}^N y^n \mathbf{w}^T \mathbf{x}^n$
- ☐ c. $-\sum_{n=1, a^n y^n=-1}^N y^n \mathbf{w}^T \mathbf{x}^n$ where $a^n = \text{sign}(\mathbf{w}^T \mathbf{x}^n)$
- ☐ d. $-\sum_{n=1}^N y^n \mathbf{w}^T \mathbf{x}^n$
- ☐ e. $-\sum_{n=1, a^n y^n=1}^N y^n \mathbf{w}^T \mathbf{x}^n$ where $a^n = \text{sign}(\mathbf{w}^T \mathbf{x}^n)$

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:26

Question 4

Not yet
answered

Marked out of
3.00

 Flag question

Under which assumption does the optimization procedure for the Perceptron converge?

Select one or more:

- ☐ a. It converges for all problems as long as the data is normalized
- ☐ b. It is never guaranteed to converge
- ☐ c. It converges only if the data is in range $[0, 1]$ and linearly separable
- ☐ d. It converges for linearly separable problems after finite many steps
- ☐ e. It converges for all problems after finite many steps

Previous page

Next page

 Slides

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 5

Not yet
answered

Marked out of
4.00

🚩 Flag question

Write down the recursion formula for delta-errors of the backpropagation algorithm using the notation of the lecture.

$$\delta^{l-1} = \dots$$

Select one:

- ☐ a. $\delta^{l-1} = \delta^l \mathbf{W}^{l-1} \text{diag}(f'(s^{l-1})) f(s^{l-1})$
- ☐ b. $\delta^{l-1} = \delta^l \mathbf{W}^{l-1} \text{diag}(f'(s^{l-1})) s^{l-1}$
- ☐ c. $\delta^{l-1} = \delta^l \mathbf{W}^l \text{diag}(f'(s^{l-1}))$
- ☐ d. $\delta^{l-1} = \delta^l \mathbf{W}^l \text{diag}(f'(s^l))$
- ☐ e. $\delta^{l-1} = \delta^{l-1} \mathbf{W}^{l-1} \text{diag}(f'(s^{l-1}))$

Time left 1:29:18

Previous page

Next page

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:11

Question 6

Not yet answered

Marked out of 2.00

 [Flag question](#)

Mark all activation functions that map all inputs to strictly **positive** values!

Select one or more:

- ☐ a. Linear
- ☐ b. SELU
- ☐ c. ReLU
- ☐ d. Sigmoid
- ☐ e. Tanh

[Previous page](#)

[Next page](#)

[◀ Slides](#)



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:06

Question 7

Not yet answered

Marked out of 2.00

 [Flag question](#)

Mark all activation functions whose derivative is not continuous!

Select one or more:

- ☐ a. Tanh
- ☐ b. SeLU
- ☐ c. Sigmoid
- ☐ d. Linear
- ☐ e. ReLU

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:29:02

Question 8

Not yet
answered

Marked out of
2.00

 [Flag question](#)

Mark all activation functions that are strictly monotonically increasing!

Select one or more:

- ☐ a. Linear
- ☐ b. Sigmoid
- ☐ c. ReLU
- ☐ d. Tanh
- ☐ e. SeLU

Previous page

Next page

 [Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:28:56

Question 9

Not yet
answered

Marked out of
2.00

 [Flag question](#)

Mark all activation functions that are bounded from above and below!

Select one or more:

- ☐ a. Tanh
- ☐ b. Sigmoid
- ☐ c. Linear
- ☐ d. SELU
- ☐ e. ReLU

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 10

Not yet answered

Marked out of 15.00

[Flag question](#)

Forward pass:

prediction: $\hat{y} =$

output layer:

hidden layer:

input layer:

inputs:

 $x = (0.5, 0.5)$ Activation (linear)
Preactivation s_3 : Activation (sigmoid) a_2 :
Preactivation s_2 : Activation (sigmoid) a_1 :
Preactivation s_1 :

$$\text{sigmoid}(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma(x)' = \sigma(x)(1 - \sigma(x))$$

Backward pass:

loss function: $L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$ actual labels: $y = 0$ prediction: $\hat{y} =$ (copy from forward pass)

output layer:

hidden layer:

input layer:

inputs:

 $x = (0.5, 0.5)$

$$\frac{\partial L}{\partial w_{1,1}^{[1]}} =$$

The preactivations of the hidden layer are denoted as s_1 , and s_2 from left to right. The respective activations, are denoted as $a_1 = \sigma(s_1)$ and identical for s_2 . The preactivation of the output layer is denoted as s_3 and the output as $\hat{y}$. The delta at the output is denoted as δ_1 , and the hidden deltas as δ_2 , and δ_3 from left to right, respectively.

The gradient of the loss w.r.t. the first weight of the first layer $\frac{\partial L}{\partial w_{1,1}^{[1]}}$ is the last to calculate.

Your answer must be the output $\hat{y}$, the delta error δ_2 as well as the gradient of the loss w.r.t. the first weight of the first layer $\frac{\partial L}{\partial w_{1,1}^{[1]}}$, therefore the answer will be (without ")

" $\hat{y}; \delta_2; \frac{\partial L}{\partial w_{1,1}^{[1]}}$ "

To give an example, your answer would look like (without ") : "1;-0.1;0.02".

Round to **3 decimal places** if necessary.

You must **separate the answers with semicolons, use dots as commas and not add unnecessary zeros**.

The following answers would be **incorrect**:

Trailing zeros: "1.0;-0.10;0.020"

Commas instead of dots: "1;-0,1;0,02"

Answer:

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:28:31

Question 11

Not yet
answered

Marked out of
3.00

 [Flag question](#)

The generalization error which is the expected loss on future data (also called risk R_{gen}) is:

Select one:

- ☐ a. A functional, i.e. an operator which maps functions to scalars
- ☐ b. A function that maps from $\mathbb{R}^N$ to $\mathbb{R}^N$
- ☐ c. A linear mapping

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:28:19

Question 12

Not yet
answered

Marked out of
3.00

 [Flag question](#)

A special method to empirically estimate the generalization error is cross-validation. Mark the true statements about it.

Select one or more:

- ☐ a. Cross-validation should also be used on huge datasets, due to lower variance of the estimator as training data is overlapping
- ☐ b. Leave-one-out cross-validation is a special case, where $N = L$, therefore each fold contains only one element.
- ☐ c. We divide the dataset in L folds, train on one folds and validate on the other folds.
- ☐ d. Cross-validation is used when we have only few datapoints and can not afford to keep a separate test set
- ☐ e. Cross-validation provides an almost unbiased estimate for the generalization error.
- ☐ f. The cross-validation error is the highest error encountered for all L different folds.

Previous page

Next page

 [Slides](#)

Jump to...



1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 13

Not yet
answeredMarked out of
4.00

Flag question

Perform the 2D Convolution operation $*$ with a **zero-padding of 1** using the following image $\mathbf{x}$ and kernel $\mathbf{w}$:

$$\mathbf{x} = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 1 & 1 \end{pmatrix}$$

and

$$\mathbf{w} = \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix}$$

Compute $\mathbf{w} * \mathbf{x}$!

Select one:

☐ a. $\begin{pmatrix} 3 & 4 & 3 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{pmatrix}$

☐ b. $\begin{pmatrix} 4 & 3 & 4 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{pmatrix}$

☐ c. $\begin{pmatrix} 2 & 2 & 3 \\ 2 & 4 & 3 \\ 1 & 2 & 2 \end{pmatrix}$

☐ d. $\begin{pmatrix} 4 \end{pmatrix}$

☐ e. $\begin{pmatrix} 2 & 2 & 2 \\ 2 & 4 & 3 \end{pmatrix}$

☐ c.
$$\begin{pmatrix} 2 & 2 & 3 \\ 2 & 4 & 3 \\ 1 & 2 & 2 \end{pmatrix}$$

☐ d.
$$(4)$$

☐ e.
$$\begin{pmatrix} 2 & 2 & 2 \\ 2 & 4 & 3 \\ 1 & 2 & 2 \end{pmatrix}$$

☐ f.
$$\begin{pmatrix} 4 & 3 & 4 \\ 3 & 4 & 3 \\ 2 & 3 & 4 \end{pmatrix}$$

☐ g.
$$\begin{pmatrix} 3 & 3 & 4 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{pmatrix}$$

☐ h.
$$\begin{pmatrix} 2 & 2 \\ 3 & 4 \end{pmatrix}$$

☐ i.
$$\begin{pmatrix} 4 & 3 & 4 \\ 2 & 4 & 3 \\ 3 & 3 & 4 \end{pmatrix}$$

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...


Question 14

Not yet answered

Marked out of 13.00

Flag question

Time left 1:27:48



Layer description	Data dimension [width, height, channel] or [#features]	Number of Parameters
Input layer	[30, 30, 3]	0
C1: 3x3 convolutions with 4 kernels, stride 3		
C2: 4x4 convolutions with 8 kernels, stride 1		
C3: 3x3 convolutions with 8 kernels, stride 1		
C4: 3x3 convolutions with 16 kernels, stride 1		
F5: Fully connected layer with 100 units		
F6: Fully connected layer with 50 units		
O7: Output layer, fully connected with 10 units		

The network above was especially created for your exam, to classify your ability to calculate the data dimensions after convolutions are applied and the number of parameters involved. It could operate on the upscaled MNIST dataset as well though. **No padding** was used in any convolutional layer. Also note that there are **no bias units** throughout the network! Images may not reflect real proportions, focus on the values rather.

Compute the dimensions after the convolutions (thus after the last convolutional layer C4) and the total number of parameters in the whole network!

The answer should be in format "[width,height,#features],#parameters" or "[#features],#parameters"!

Example answer: "[64,64,3],0" (without the "-marks")!

Answer:

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:27:37

Question 15

Not yet
answered

Marked out of
3.00

 [Flag question](#)

What is the problem with Dropout concerning training-mode and test mode?

Select one:

- ☐ a. There is no problem concerning dropout concerning training and test mode.
- ☐ b. During test mode the same neurons are dropped out as during training.
- ☐ c. During testing the outputs of the backward pass need to be rescaled by the the probability that a unit activation is not dropped.
- ☐ d. During testing the outputs of the backward pass need to be rescaled by the inverse dropout probability.
- ☐ e. During testing there is a bias in the expectation of the pre-activations compared to the expectation during training with dropout.

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 16

Not yet answered

Marked out of 10.00

[Flag question](#)

Assume you want to develop an AI for playing a computer game. The computer game involves **green tiles** (see green parts on the image below) on a screen (white part on the image below) that have to be moved to a certain place.

You chose a convolutional neural network $g(\mathbf{x}; \mathbf{w})$ as one component of your AI to recognize whether the **green tile** is on the image $x \in \mathbb{R}^{A \times B \times C}$ or not. Images have height A , width B and C channels. For this convolutional network you want to guarantee that it outputs the same prediction for all possible 90-degree-rotations of the tile and all possible positions of the tile. The usual convolution operation (stride 1, padding 1) for multiple kernels is used $\mathbf{h} = \text{Conv}(\{w_1, w_2, w_3, w_4\}, \mathbf{x})$. Each of the four kernels is convolved with the image $w_i * \mathbf{x}$, and the resulting feature maps are stacked so that $\mathbf{h} \in \mathbb{R}^{A \times B \times 4}$. We denote an element of this quantity as $h_{a,b,c}$.

In the image below, you find an example of the positive class of images. The negative class consists of images that do not contain the green tile.

Image 1						Image 2					
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	0	0	0	1	0	0	0
0	0	1	1	0	0	0	0	1	1	0	0
0	0	0	0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

Pick a set of kernels (columns a, b, c, d, e, f in the image below) and neural network that **guarantee** that the neural network predicts **the same value** for **all possible locations** and **90degree-rotations** of the green tile! Note that multiple answers can be correct!

	a	b	c	d	e	f
w1	0 1 0 0 1 1 0 0 1	1 0 1 0 1 0 0 0 0	0 1 0 0 1 1 0 1 0	0 0 0 0 1 1 1 1 0	0 1 0 0 1 0 0 1 0	1 0 0 1 1 1 0 0 0
w2	1 0 0 0 1 0 1 0 1	0 0 1 0 1 0 0 0 1	0 0 1 1 1 1 0 0 0	0 1 0 0 1 1 0 0 1	0 0 0 1 1 1 0 0 0	0 1 1 1 1 0 0 0 0
w3	0 1 0 0 1 0 1 0 1	0 0 0 0 1 0 1 0 1	1 0 0 0 1 0 0 1 0	0 1 1 1 1 0 0 0 0	0 0 1 0 1 0 1 0 0	0 1 1 0 1 0 1 0 0
w4	0 0 1 1 1 1 0 0 0	1 0 0 0 1 0 1 0 0	1 1 0 0 1 0 0 0 1	1 0 0 1 1 0 0 1 0	1 0 0 0 1 0 0 0 1	0 0 1 1 1 1 0 0 0

Select one or more:

☐ 1. Kernel set: **a** and neural network $g: \hat{y} = \sigma(\frac{1}{AB} \sum_{a=0}^{A-1} \sum_{b=0}^{B-1} (\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen.

☐ 2. Kernel set: **b** and neural network $g: \hat{y} = \sigma(\max_{a,b} (\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen and

	a	b	c	d	e	f
w1	0 1 0 0 1 1 0 0 1	1 0 1 0 1 0 0 0 0	0 1 0 0 1 1 0 1 0	0 0 0 0 1 1 1 1 0	0 1 0 0 1 0 0 1 0	1 0 0 1 1 1 0 0 0
w2	1 0 0 0 1 0 1 0 1	0 0 1 0 1 0 0 0 1	0 0 1 1 1 1 0 0 0	0 1 0 0 1 1 0 0 1	0 0 0 1 1 1 0 0 0	0 1 1 1 1 0 0 0 0
w3	0 1 0 0 1 0 1 0 1	0 0 0 0 1 0 1 0 1	1 0 0 0 1 0 0 1 0	0 1 1 1 1 0 0 0 0	0 0 1 0 1 0 1 0 0	0 1 1 0 1 0 1 0 0
w4	0 0 1 1 1 1 0 0 0	1 0 0 0 1 0 1 0 0	1 1 0 0 1 0 0 0 1	1 0 0 1 1 0 0 1 0	1 0 0 0 1 0 0 0 1	0 0 1 1 1 1 0 0 0

Select one or more:

- ☐ 1. Kernel set: **a** and neural network $g: \hat{y} = \sigma(\frac{1}{AB} \sum_{a=0}^{A-1} \sum_{b=0}^{B-1} (\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen.
- ☐ 2. Kernel set: **b** and neural network $g: \hat{y} = \sigma(\max_{a,b}(\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen and $\max_{a,b}$ indicates the maximal value of all positions.
- ☐ 3. Kernel set: **c** and neural network $g: \hat{y} = \sigma(\frac{1}{AB} \sum_{a=0}^{A-1} \sum_{b=0}^{B-1} (\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen.
- ☐ 4. Kernel set: **d** and neural network $g: \hat{y} = \sigma(\frac{1}{AB} \sum_{a=0}^{A-1} \sum_{b=0}^{B-1} (\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen.
- ☐ 5. Kernel set: **f** and neural network $g: \hat{y} = \sigma(\max_{a,b}(\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen and $\max_{a,b}$ indicates the maximal value of all positions.
- ☐ 6. Kernel set: **e** and neural network $g: \hat{y} = \sigma(\max_{a,b}(\max_c h_{a,b,c}))$, where $\max_c$ indicates that the maximal value of all featuremaps/channels is chosen and $\max_{a,b}$ indicates the maximal value of all positions.

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question 17

Not yet
answered

Marked out of
4.00

Flag question

AdaGrad adapts the step size by remembering past gradients $\mathbf{g}^{(t)}$. The parameter wise learning rates are then scaled down for parameters that already experienced large gradients and up for parameters that experienced low gradients so far. For this end, we store the sum of squared past gradients $\mathbf{v}^{(t)} = \sum_{s=1}^t (\mathbf{g}^{(s)})^2$, where the squared is computed component wise. We use this to update the parameter vector $\mathbf{w}$ similar to the basic gradient descent algorithm. Select the correct formula that AdaGrad uses for this update.

Time left 1:26:59

Select one:

- ☐ a.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} + \eta \frac{1}{\sqrt{\mathbf{v}^{(r)}} + \epsilon} \odot \mathbf{g}^{(r)}$$
- ☐ b.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} + \eta \frac{1}{\sqrt{\mathbf{v}^{(r-1)}} + \epsilon} \odot \mathbf{g}^{(r)}$$
- ☐ c.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \eta \frac{1}{\sqrt{\mathbf{v}^{(r-1)}} + \epsilon} \odot \mathbf{g}^{(r)}$$
- ☐ d.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} + \eta \frac{1}{\sqrt{\mathbf{v}^{(r)}} + \epsilon} \cdot \mathbf{g}^{(r)}$$
- ☐ e.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \eta \frac{1}{\sqrt{\mathbf{v}^{(r-1)}} + \epsilon} \cdot \mathbf{g}^{(r)}$$
- ☐ f.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \eta \frac{1}{\sqrt{\mathbf{v}^{(r)}} + \epsilon} \cdot \mathbf{g}^{(r)}$$
- ☐ g.
$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \eta \frac{1}{\sqrt{\mathbf{v}^{(r)}} + \epsilon} \odot \mathbf{g}^{(r)}$$

Previous page

Next page

Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:26:54

Question **18**

Not yet
answered

Marked out of
2.00

 [Flag question](#)

The loss surface of neural networks (loss considered as the function of weights) is in general strictly convex, true or false?

Select one:

- ☐ a. True, because neural networks can approximate any function arbitrarily close according to the universal function approximator theorem.
- ☐ b. False, the convexity of the loss function depends on the choice of the optimizer
- ☐ c. False, only Sigmoid or Tanh activation leads to strictly convex loss surfaces
- ☐ d. True, due to the stacking of linear layers and activation functions the loss surface is strictly convex
- ☐ e. False, due to stacking of layers the loss surface increases in complexity

[Previous page](#)

[Next page](#)

[◀ Slides](#)



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:26:47

Question 19

Not yet answered

Marked out of 6.00

[Flag question](#)

Suppose we are performing gradient descent on a quadratic objective: $R(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w}$ with symmetric matrix $\mathbf{A}$. The dynamics of gradient descent with learning rate η can be analyzed in terms of the spectral decomposition $\mathbf{A} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T$, where $\mathbf{Q}$ is an orthogonal matrix containing the eigenvectors of $\mathbf{A}$, and $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_D)$ is a diagonal matrix containing the eigenvalues of $\mathbf{A}$ in descending order, where $\lambda_1 = \lambda_{\max}$ is the largest eigenvalue. Then we have

$$\mathbf{w}^t = \mathbf{Q}(\mathbf{I} - \eta \mathbf{\Lambda})^t \mathbf{Q}^T \mathbf{w}^0.$$

Based on this formula, what is the value of C such that the gradient descent iterates diverge for $\eta > C$ but converge for $\eta < C$?

Select one:

- ☐ a. $C = \frac{1}{D}$
- ☐ b. $C = \lambda_{\max}$
- ☐ c. $C = \frac{2}{\lambda_{\max}}$
- ☐ d. $C = 1$
- ☐ e. $C = \frac{1}{\lambda_{\max}}$

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:26:39

Question 20

Not yet answered

Marked out of 6.00

[Flag question](#)

Assume you have a convolutional layer of 4x4 filters with a total of 16 filters and stride 2 that is applied to inputs of dimensions [32,32,3] (height, width, features) and you want to initialize the weights of this conv-layer according to **Glorot's Initialization** scheme. Further assume that you want to use a normal distribution. What are the correct parameters for this distribution?

- ☐ a. $W_{ij} \sim \mathcal{N}(0, \frac{2}{4 \times 4 \times 16})$
- ☐ b. $W_{ij} \sim \mathcal{N}(0, \frac{2}{4 \times 4 \times 3 + 4 \times 4 \times 16})$
- ☐ c. $W_{ij} \sim \mathcal{N}(0, \frac{1}{4 \times 4 \times 3})$
- ☐ d. $W_{ij} \sim \mathcal{N}(0, \frac{2}{32 \times 32 \times 3 + 32 \times 32 \times 16})$
- ☐ e. $W_{ij} \sim \mathcal{N}(1, \frac{2}{4 \times 4 \times 16})$

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Time left 1:26:33

Question **21**

Not yet
answered

Marked out of
2.00

 [Flag question](#)

Inputs for a feed-forward neural network should be scaled to the interval $[0, 100]$ - true or false?

Select one:

- ☐ a. True, because large numbers leads to larger gradients and therefore faster learning
- ☐ b. True, because large positive numbers counteract the vanishing gradient problem
- ☐ c. True, because gradient descent only works with positive numbers
- ☐ d. False, because large numbers leads to larger gradients and therefore unstable learning
- ☐ e. False, because gradient descent only works with negative numbers

[Previous page](#)

[Next page](#)

[◀ Slides](#)

Jump to...



Quiz navigation

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22		

Finish attempt ...

Question **22**

Not yet
answered

Marked out of
4.00

 [Flag question](#)

[Previous page](#)

 [Slides](#)

Time left 1:26:26

Which of the following statements about Batch Normalization are true?

Select one or more:

- ☐ a. Each neuron learns a mean batch activation μ and standard deviation σ .
- ☐ b. Batch normalization introduces two learnable parameters β and γ for each neuron, which can scale and shift the normalized activation.
- ☐ c. In the standard formulation presented in the lecture, batch normalization aims to standardize the neuron pre-activations.
- ☐ d. Batch normalization introduces two learnable parameters β and γ for each layer, which can scale and shift the normalized activation.
- ☐ e. The network learns a single batch mean activation μ and standard deviation σ that are the same for all neurons in a layer.

Finish attempt ...

Jump to...

