

# 1 Deep Reinforcement Learning Algorithms

| Algorithm            | Family          | Main Idea                    | Core Equation                                                                                   | Policy    | Model   |
|----------------------|-----------------|------------------------------|-------------------------------------------------------------------------------------------------|-----------|---------|
| Greedy               | Bandits         | Always choose best action    | $A_t = \arg \max_a Q_t(a)$                                                                      | On        | Free    |
| $\epsilon$ -Greedy   | Bandits         | Random exploration           | $A_t = \begin{cases} \arg \max_a Q(a), & 1 - \epsilon \\ \text{random}, & \epsilon \end{cases}$ | On        | Free    |
| UCB                  | Bandits         | Optimism under uncertainty   | $A_t = \arg \max_a \left( Q(a) + c \sqrt{\frac{\log t}{N(a)}} \right)$                          | On        | Free    |
| Thompson Sampling    | Bandits         | Sample from posterior        | $a = \arg \max_a \tilde{q}(a), \quad \tilde{q}(a) \sim P(q(a))$                                 | On        | Free    |
| Gradient Bandits     | Policy-based    | Learn action preferences     | $\pi(a) = \frac{e^{H(a)}}{\sum_b e^{H(b)}}$                                                     | On        | Free    |
| Dynamic Programming  | Value-based     | Solve known MDP              | Bellman equations                                                                               | Either    | Based   |
| Policy Iteration     | DP              | Evaluate then improve policy | $\pi'(s) = \arg \max_a q_\pi(s, a)$                                                             | Either    | Based   |
| Value Iteration      | DP              | Direct Bellman optimization  | $V(s) \leftarrow \max_a [r + \gamma V(s')]$                                                     | Either    | Based   |
| Monte Carlo          | Value-based     | Learn from episodes          | $V(s) \leftarrow V(s) + \alpha(G_t - V(s))$                                                     | On        | Free    |
| Temporal Difference  | Value-based     | Bootstrap estimates          | $V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$                                        | On        | Free    |
| SARSA                | Value-based     | Use executed actions         | $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$                           | On        | Free    |
| Q-Learning           | Value-based     | Learn optimal values         | $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$                 | Off       | Free    |
| DQN                  | Deep RL         | Neural network Q-function    | $L = (y - Q(s, a; \theta))^2$                                                                   | Off       | Free    |
| REINFORCE            | Policy Gradient | Direct policy optimization   | $\theta \leftarrow \theta + \alpha G_t \nabla_\theta \log \pi_\theta(a s)$                      | On        | Free    |
| Actor-Critic         | Hybrid          | Actor acts, critic evaluates | $\delta = r + \gamma V(s') - V(s)$                                                              | On        | Free    |
| A2C/A3C              | Actor-Critic    | Use advantage estimates      | $A(s, a) = Q(s, a) - V(s)$                                                                      | On        | Free    |
| TRPO                 | Policy Gradient | Constrained updates          | $\max E[r_t A_t] \quad s.t. \quad D_{KL} \leq \delta$                                           | On        | Free    |
| PPO                  | Policy Gradient | Clipped optimization         | $L = \min(r_t A_t, \text{clip}(r_t, 1 - \epsilon, 1 + \epsilon) A_t)$                           | On        | Free    |
| DDPG                 | Actor-Critic    | Continuous actions           | $\nabla_\theta J = \nabla_a Q(s, a) \nabla_\theta \mu_\theta(s)$                                | Off       | Free    |
| TD3                  | Actor-Critic    | Twin critics                 | $y = r + \gamma \min(Q_1, Q_2)$                                                                 | Off       | Free    |
| SAC                  | Actor-Critic    | Reward + entropy             | $J = E[R] + \alpha H(\pi)$                                                                      | Off       | Free    |
| Dyna-Q               | Hybrid          | Model + simulation           | Q-learning + planning                                                                           | Off       | Based   |
| Prioritized Sweeping | Planning        | Update important states      | Bellman error priority                                                                          | Off       | Based   |
| AlphaZero            | Planning        | Self-play + MCTS             | $U(s, a) = Q(s, a) + c P(s, a) \frac{\sqrt{N}}{1+n(s, a)}$                                      | Self-play | Based   |
| MuZero               | Planning        | Learn world model            | Learn networks $(h, g, f)$                                                                      | Self-play | Learned |
| Dreamer              | World Models    | Imagine trajectories         | Optimize imagined returns                                                                       | Either    | Based   |
| MPC                  | Planning        | Replan every step            | $\max_{a_{0:H}} \sum_t r_t$                                                                     | N/A       | Based   |
| Random Shooting      | MPC             | Random trajectories          | $\arg \max_i R_i$                                                                               | N/A       | Based   |
| CEM                  | Optimization    | Fit elite samples            | Gaussian update                                                                                 | N/A       | Based   |
| iLQR                 | Control         | Linearize dynamics           | Riccati recursion                                                                               | N/A       | Based   |

## 2 Fundamental Reinforcement Learning Equations

### 2.1 Return

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \tag{1}$$

The return represents the discounted cumulative future reward.

### 2.2 Bellman Expectation Equation

$$v_{\pi}(s) = \sum_a \pi(a|s) \left[ r(s,a) + \gamma \sum_{s'} P(s'|s,a) v_{\pi}(s') \right] \tag{2}$$

### 2.3 Bellman Optimality Equation

$$v^*(s) = \max_a \left[ r(s,a) + \gamma \sum_{s'} P(s'|s,a) v^*(s') \right] \tag{3}$$

$$q^*(s,a) = r + \gamma \sum_{s'} P(s'|s,a) \max_{a'} q^*(s',a') \tag{4}$$

## 3 RL Family Tree

```
Reinforcement Learning
|
|-- Value-Based
|   |-- TD
|   |-- SARSA
|   |-- Q-Learning
|   |-- DQN
|
|-- Policy-Based
|   |-- REINFORCE
|   |-- TRPO
|   |-- PPO
|
|-- Actor-Critic
|   |-- A2C/A3C
|   |-- DDPG
|   |-- TD3
|   |-- SAC
|
|-- Model-Based
|   |-- Dyna-Q
|   |-- MPC
|   |-- AlphaZero
|   |-- MuZero
|   |-- Dreamer
```