

DRL Exam Cheat Sheet

Deep Reinforcement Learning (SS26)

(c) Tobias Nolz

1. Basics & MDPs

MDP: Tuple (S, A, P, R, γ) . S : States, A : Actions, $P(s'|s, a)$: Trans. dynamics, $R(s, a) = \mathbb{E}[R|s, a]$: Reward, $\gamma \in [0, 1]$.

Markov Property: $\mathbb{P}(S_{t+1}, R_{t+1} | S_t, A_t, \dots) = \mathbb{P}(S_{t+1}, R_{t+1} | S_t, A_t)$.

Return: $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$.

Policy: $\pi(a|s) \in [0, 1]$, $\sum_a \pi(a|s) = 1$. Goal: $\max_{\pi} \mathbb{E}_{\pi}[G_t]$.

Value Funcs: $V^{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$, $Q^{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$.

Relation: $V^{\pi}(s) = \sum_a \pi(a|s) Q^{\pi}(s, a)$.

Optimal: $V^*(s) = \max_{\pi} V^{\pi}(s)$, $Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a)$.

Advantage: $A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$.

Bellman Expectation Equations:

$$V^{\pi}(s) = \sum_a \pi(a|s) \left[r(s, a) + \gamma \sum_{s'} P(s'|s, a) V^{\pi}(s') \right]$$

$$Q^{\pi}(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \sum_{a'} \pi(a'|s') Q^{\pi}(s', a')$$

Bellman Optimality Equations:

$$V^*(s) = \max_a \left[r(s, a) + \gamma \sum_{s'} P(s'|s, a) V^*(s') \right]$$
$$Q^*(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^*(s', a')$$

2. Multi-Armed Bandits

$q(a) = \mathbb{E}[R_t | A_t = a]$. Sample avg: $Q_t(a) = \frac{1}{N_t(a)} \sum_{\tau \leq t} R_{\tau} \mathbf{1}[A_{\tau} = a]$.

Incremental: $Q_{t+1}(a) = Q_t(a) + \frac{1}{N_t(a)} [R_t - Q_t(a)]$.

ϵ -greedy: $A_t = \arg \max_a Q_t(a)$ w.p. $1 - \epsilon$.

UCB: $A_t = \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$.

Regret: $L_T = \sum_{t=1}^T [q^* - q(A_t)]$. ϵ -greedy gives linear regret. UCB yields logarithmic regret $O(\ln T)$.

3. Dynamic Programming

Requires known P, R . Based on sweeping Bellman equations. Prediction/policy evaluation (linear $\Rightarrow$ closed form $O(|S|^3)$). Control/policy optimization (nonlinear $\Rightarrow$ iterative).

Policy Eval: $V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s', r} P[r + \gamma V_k(s')]$.

Policy Improv. Theorem: If $Q^{\pi}(s, \pi'(s)) \geq V^{\pi}(s) \forall s$, then $V^{\pi'}(s) \geq V^{\pi}(s) \forall s$. Greedy: $\pi'(s) = \arg \max_a Q^{\pi}(s, a)$.

Policy Iteration: Evaluate V^{π} , then greedy improvement. Repeat. Guaranteed to converge.

Value Iteration: Combine eval/improve in one sweep: $V_{k+1}(s) = \max_a \sum_{s', r} P[r + \gamma V_k(s')]$. Extract optimal policy at the end.

4. Tabular Model-Free Prediction

MC: Full episodes. Unbiased, high variance. $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$.

TD(0): 1-step bootstrap. Biased, lower variance. Online.

$$V(S_t) \leftarrow V(S_t) + \alpha \underbrace{(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))}_{\delta_t \text{ (TD error)}}$$

n -step TD: $G_t^{(n)} = \sum_{k=1}^n \gamma^{k-1} R_{t+k} + \gamma^n V(S_{t+n})$.

Update: $V(S_t) \leftarrow V(S_t) + \alpha(G_t^{(n)} - V(S_t))$.

Eligibility Traces TD(λ): Forward: $G_t^{\lambda} = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$. ($\lambda=0 \Rightarrow$ TD(0), $\lambda=1 \Rightarrow$ MC.)

Backward: $z_t = \gamma \lambda z_{t-1} + \nabla V(S_t)$; $\Delta w = \alpha \delta_t z_t$.

5. Tabular Model-Free Control

SARSA (On-Policy TD): $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$, ϵ -greedy for behavior + update. **GLIE:** $\epsilon_t \rightarrow 0$, $\sum_t \epsilon_t = \infty \Rightarrow$ converges to q^* .

Expected SARSA: $Q \leftarrow Q + \alpha[R + \gamma \sum_a \pi(a|S') Q(S', a) - Q]$. Lower variance.

Q-Learning (Off-Policy TD): $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') - Q(S_t, A_t)]$ Converges to q^* : Robbins-Monro ($\sum \alpha_t = \infty$, $\sum \alpha_t^2 < \infty$)

Double Q-Learning: Decouples selection/evaluation to fix overestimation bias:

$Q^A \leftarrow Q^A + \alpha[R + \gamma Q^B(S', \arg \max_a Q^A(S', a)) - Q^A]$.

Importance Sampling: $\rho_t = \pi(A_t | S_t) / \mu(A_t | S_t)$. Full traj: $\rho_t^T = \prod_{k=t}^{T-1} \rho_k$.

6. Value Function Approx & DQN

Scale to large state spaces. Approximate $Q_w(s, a)$ or $V_w(s)$.

Linear FA: $V_w(s) = w^T x(s)$, update $\Delta w = \alpha \delta_t x(s)$. Unique global minimum.

Gradient Descent: MC: $\Delta w = \alpha(G_t - V_w(S_t)) \nabla_w V_w(S_t)$. Linear MC $\rightarrow$ global opt.

Semi-Gradient Methods: Target depends on w , so not true gradient. TD(0): $\Delta w = \alpha \delta_t \nabla_w V_w(S_t)$. TD(λ): $\Delta w = \alpha(G_t^{\lambda} - V_w(S)) \nabla_w V_w(S)$. Linear TD converges, non-linear can diverge.

DQN Objective: $L(w) = \mathbb{E}_{\mathcal{D}}[(Y - Q_w(s, a))^2]$

Target $Y = r + \gamma \max_{a'} Q_{w^-}(s', a')$.

DQN Core (two stabilization tricks):

- Experience Replay:** Random mini-batches from $\mathcal{D}$. Breaks temporal correlation, reuses data.
- Fixed Target Net (w^-):** $w^- \leftarrow \tau w + (1 - \tau) w^-$. Stabilizes targets.

Rainbow (6 extensions on top of DQN):

- Double DQN:** $Y = r + \gamma Q_{w^-}(s', \arg \max_a Q_w(s', a))$. Fixes overestimation.
- Dueling Networks:** $Q(s, a) = V(s) + A(s, a) - \frac{1}{|A|} \sum_{a'} A(s, a')$.
- Prioritized Replay (PER):** $p_i \propto |\delta_i|^w$. Bias corrected via IS weights $w_i = (N p_i)^{-\beta}$.
- Distributional (C51):** Predict categorical distribution $Z(s, a)$ of returns.
- Multi-Step Returns:** n -step targets reduce bias, faster reward propagation.
- Noisy Nets:** $y = (\mu^w + \sigma^w \odot \epsilon^w) x + (\mu^b + \sigma^b \odot \epsilon^w)$. State-dep. exploration.

7. Policy Gradients

Optimize $\pi_{\theta}(a|s)$ directly. Handles continuous/high-dim actions, enables stochastic policies.

Policy Objectives:

Episodic return (d_0 : initial-state distribution):

$$J(\theta) = \mathbb{E}_{S_0 \sim d_0, \pi_{\theta}}[G_0] = \mathbb{E}_{S_0 \sim d_0} \left[\sum_a \pi_{\theta}(a|S_0) Q^{\pi_{\theta}}(S_0, a) \right]$$

Average reward (continuing tasks; $d^{\pi_{\theta}}$: stationary dist.):

$$J(\theta) = \mathbb{E}_{\pi_{\theta}}[R_{t+1}] = \sum_s d_{\pi_{\theta}}(s) \sum_a \pi_{\theta}(a|s) r(s, a)$$

Log-Derivative Trick: $\nabla_{\theta} \pi_{\theta}(a|s) = \pi_{\theta}(a|s) \nabla_{\theta} \log \pi_{\theta}(a|s)$.

Policy Gradient Theorem (env. dynamics $p(s'|s, a)$ cancel out):

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi}(s, a)]$$

Policy Gradient Learning Signals Ψ_t :

General parameter update:

$$\theta \leftarrow \theta + \alpha \gamma^t \Psi_t \nabla_{\theta} \log \pi_{\theta}(A_t | S_t)$$

where the learning signal Ψ_t can be substituted with:

$$\begin{array}{ll} G_t & \text{REINFORCE (MC return, high variance)} \\ G_t - V_w(S_t) & \text{REINFORCE w/ baseline } b(s) \\ \delta_t = R_{t+1} + \gamma V_w(S_{t+1}) - V_w(S_t) & \text{1-Step AC (TD(0))} \\ G_t^{(n)} - V_w(S_t) & \text{n-Step AC} \\ \hat{A}_t^{\text{GAE}(\lambda)} = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l} & \text{TD}(\lambda) / \text{GAE (PPO)} \\ Q_w(S_t, A_t) & \text{Q-based AC (DDPG, SAC)} \end{array}$$

8. Actor-Critic & Stable Updates

Actor updates π_{θ} ; Critic estimates V_w or Q_w (via $\Delta w = \alpha w \Psi_t \nabla_w V_w$).

GAE Trade-offs: $\lambda=0 \Rightarrow$ TD/high bias; $\lambda=1 \Rightarrow$ MC/high variance.

KL Div.: $D_{\text{KL}}(\pi_{\text{old}} \| \pi_{\theta}) = \mathbb{E}_s [\sum_a \pi_{\text{old}} \log(\pi_{\text{old}} / \pi_{\theta})]$.

NPG: $D_{\text{KL}} \approx \frac{1}{2} \delta^T F \delta$ (Fisher matrix F). $\Delta \theta = \alpha F^{-1} \nabla J$. Expensive (2nd order). Only local approx.

TRPO: $\max_{\theta} \mathbb{E} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\text{old}}(a|s)} \hat{A} \right]$ s.t. $D_{\text{KL}}(\pi_{\text{old}} \| \pi_{\theta}) \leq \epsilon$.

Conj. grad., second order. Small, reliable policy updates.

PPO: $r_t(\theta) = \pi_{\theta} / \pi_{\text{old}}$. First-order, multiple epochs per batch. Limit destructive updates.

$$L^{\text{CLIP}} = \mathbb{E} \left[\min \left(r_t \hat{A}_t, \text{clip}(r_t, 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$$

Combined: $L^{\text{PPO}} = L^{\text{CLIP}} - c_1 L^{\text{VF}} + c_2 S[\pi_{\theta}]$ (entropy bonus).

9. Continuous Control

Gaussian Policy: $A_t \sim \mathcal{N}(\mu_{\theta}(S_t), \Sigma_{\theta}(S_t))$, score: $\nabla_{\theta} \log \pi_{\theta} = \frac{A_t - \mu_{\theta}}{\sigma^2} \nabla_{\theta} \mu_{\theta}$. See REINFORCE update.

DPG: Learn $Q_w(s, a) \approx Q^{\mu}(s, a)$. Deterministic policy $A_t = \mu_{\theta}(S_t)$. Chain rule through critic:

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[\nabla_a Q_w(s, a) \Big|_{a=\mu_{\theta}(s)} \nabla_{\theta} \mu_{\theta}(s) \right]$$

DDPG (DQN+DPG): Off-policy, replay buffer, target nets w^-, θ^- . $a = \mu_{\theta}(s)$ (learned maximizer instead). **Critic:** $\min_w \mathbb{E}[(r + \gamma Q_{w^-}(s', \mu_{\theta^-}(s')) - Q_w(s, a))^2]$. **Actor:** $\max_{\theta} \mathbb{E}[Q_w(s, \mu_{\theta}(s))]$. Exploration via O-U noise. **TD3** (fixes DDPG overestimation – 3 tricks):

- Clipped Double-Q:** $Y = r + \gamma \min_{i=1,2} Q_{\theta_i^{\pm}}(s', \tilde{a})$.
- Delayed Actor Updates:** Update μ_{θ} every d critic steps.

3. **Target Policy Smoothing:** $\tilde{a} = \mu_{\theta'}(s') + \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$.

SAC (Max-Entropy, off-policy, stochastic): $\mathcal{H}(\pi) = -\mathbb{E}[\log \pi]$.

$$J(\pi) = \mathbb{E}_{\pi} \left[\sum_t \gamma^t (R_{t+1} + \alpha \mathcal{H}(\pi(\cdot | S_t))) \right]$$

Soft Bellman: $Y = r + \gamma \mathbb{E}_{a'} [\min_i Q_{w_i^-}(s', a') - \alpha \log \pi(a' | s')]$.

Actor: $\max_{\theta} \mathbb{E}_s [\mathbb{E}_{a \sim \pi_{\theta}} [Q_w(s, a)] + \alpha \mathcal{H}(\pi_{\theta}(\cdot | s))]$.

Reparam.: $a = \mu_{\theta}(s) + \sigma_{\theta}(s) \odot \epsilon$, $\epsilon \sim \mathcal{N}(0, I)$ – gradients flow to θ .

10. Exploration & Representation

Bootstrap / Thompson Sampling: K Q-heads; sample one head per episode for acting. Bootstrap masks ensure data diversity across heads.

Directed Exploration: aug. reward $r_t = r_t^{\text{ext}} + \beta r_t^{\text{int}}$.

Count-Based: $r^{\text{int}}(s) \propto 1/\sqrt{N(s)}$. Intractable in large spaces; use pseudo-counts or hash-based approximations.

ICM: Learn features ϕ via *inverse* model $\hat{a}_t = g(\phi(s_t), \phi(s_{t+1}))$ – filters dimensions the agent cannot control. Forward model then predicts next feature:

$$r_{\text{int}} = \|\mathbf{f}(\phi(s_t), a_t) - \phi(s_{t+1})\|^2$$

Vulnerable to *noisy TV* (irreducible stochastic noise).

RND: Predictor f_{θ} is trained to match *fixed* random target f_{tgt} : Deterministic target $\Rightarrow$ mitigates noisy TV.

$$r_{\text{int}} = \|\mathbf{f}_{\theta}(s') - \mathbf{f}_{\text{tgt}}(s')\|^2 \quad (\downarrow \text{ as } s' \text{ is visited more})$$

Disagreement: $r_{\text{int}} \propto \text{Var}_k[\hat{s}_{t+1}^{(k)}]$ from dynamics model ensemble; high where model is uncertain.

State Rep. Learning ($o_t \rightarrow x_t$): Goals: Structure, Continuity, Sparsity. *Bisimulation:* Latent distance $\propto$ transition & reward similarity. *Aux. Tasks:* Add reconstruction or forward/inverse dynamics prediction. *Data Aug:* Crop/jitter to encourage invariance to visual nuisances. *Contrastive:* Pull similar states (+) close, push different (–) apart ($\mathcal{L}_{\text{NCE}}$). *Attention:* Mask/patch self-attention to filter visual distractors.

11. Credit Assignment

Three root causes: *Depth:* Long delay between action & reward $\rightarrow$ recency bias. *Density:* Sparse rewards $\rightarrow$ weak or absent learning signal. *Breadth:* Many contributing actions $\rightarrow$ diluted per-action credit.

Addressing Depth: n -step returns / TD(λ) bridge short and long horizons. GAE: exponential weighting of δ_t 's, $\hat{A}_t^{\text{GAE}} = \sum_l (\gamma \lambda)^l \delta_{t+l}$. **RUDDER:** learn return predictor $\hat{G}_t$; redistribute $\tilde{r}_{t+1} = \hat{G}_{t+1} - \hat{G}_t$.

Addressing Density: Reward shaping.

Addressing Breadth: Advantage $A(s, a) = Q - V$ isolates the contribution of action a over the average.

HER: Goal-conditioned policy; relabel failed traj. $(s_0, \dots, s_T)$ with achieved goal $g' = s_T$. Every episode becomes a success for *some* goal $\Rightarrow$ dense signal in sparse-reward envs. Lim: requires a goal-conditioned MDP.

12. Imitation Learning (IL)

BC (Behavioral Cloning): Supervised MLE on expert demos:

$$\max_{\theta} \mathbb{E}_{(s,a) \sim \mathcal{D}_{\text{exp}}} [\log \pi_{\theta}(a|s)]$$

Flaw – Distribution Shift: agent deviates $\rightarrow$ visits states absent from $\mathcal{D} \rightarrow$ errors compound.

Ross & Bagnell 2010: $J(\pi^*) - J(\hat{\pi}_{\text{BC}}) \leq T^2\epsilon$ where $\epsilon = \max_s \mathcal{L}^{\pi^*}(s)$. Quadratic in horizon T .

DAgger (Dataset Aggregation): Trains on learner’s *own* distribution. Reduces bound to $O(T\epsilon)$.

1. Init $\mathcal{D} \leftarrow \mathcal{D}_{\text{exp}}$. Train $\hat{\pi}_1$ via BC.
2. **Repeat:** Roll out $\hat{\pi}_i$, collect visited states S_i .
3. Query expert: get labels $\pi^*(s) \forall s \in S_i$.
4. $\mathcal{D} \leftarrow \mathcal{D} \cup \{(S_i, \pi^*(S_i))\}$; retrain $\hat{\pi}_{i+1}$.

Other Failure Modes: *Non-Markovian expert:* Expert uses history $\Rightarrow$ Markov BC policy is confused. Fix: use RNN / memory. *Causal confusion:* Policy latches on spurious correlates (e.g. brake light, not velocity). *Multimodal actions:* MSE regressor averages modes $\rightarrow$ unsafe mean action. Fix: diffusion/flow policies or mixture density networks.

13. Offline RL

Fixed dataset $\mathcal{D} = \{(s, a, r, s')\}$ from behaviour policy π_β ; no env interaction.

Core Challenge: Bootstrapping $Q(s', a')$ for OOD actions $a' \notin \mathcal{D}$ causes overestimation that propagates via Bellman $\Rightarrow$ policy collapse.

Behavior Reg: Constrain $\pi \rightarrow \pi_\beta$ to avoid OOD errors: $\max J(\pi)$ s.t. $D(\pi \| \pi_\beta) \leq \delta$. *Methods:* **1)** KL Penalty (λD_{KL}), **2)** BC Aux Loss ($\lambda \mathcal{L}_{\text{BC}}$), **3)** Support Bound ($\pi(a|s) = 0$ if OOD).

CQL (Conservative Q-Learning): Explicitly penalise OOD Q-values while supporting in-data ones:

$$\min_Q \underbrace{\alpha(\mathbb{E}_{s,a \sim \mu}[Q] - \mathbb{E}_{s,a \sim \mathcal{D}}[Q])}_{\text{conservative penalty}} + \underbrace{\frac{1}{2}\mathcal{L}_{\text{TD}}}_{\text{Bellman error}}$$

Provably conservative: $Q_{\text{CQL}}(s, a) \leq Q^\pi(s, a) \forall (s, a)$.

IQL: Avoids OOD queries entirely. Replaces $\max_a$ with expextile regression: $V(s) = \arg \min_V \mathbb{E}_{a \sim \mathcal{D}}[\mathcal{L}^\tau(Q(s, a) - V(s))]$. *Loss:* $\mathcal{L}^\tau(u) = |\tau - 1[u < 0]|u|^2$. For $\tau \approx 1$, $V(s) \approx \max_{a \in \mathcal{D}} Q(s, a)$ (soft max over in-data actions).

Decision Transformer (DT): Return-to-go $\hat{R}_t = \sum_{t' \geq t} r_{t'}$. GPT autoregressively models $(\hat{R}_t, s_t, a_t)$ sequences; at test time condition on desired $\hat{R}_0$. *No Bellman:* trained purely supervised. *No stitching:* cannot combine sub-trajectories from $\mathcal{D}$ to find better paths (unlike CQL/IQL).

14. Inverse RL & RLHF

IRL: Recover $R_\psi(s, a)$ from expert demos $\mathcal{D}_{\text{exp}}$. Ill-posed ($R=0$ trivially optimal).

Max-Margin IRL: Expert must beat all policies by margin $D(\pi_{\text{exp}}, \pi)$:

$$\min_\psi \|\psi\|^2 \text{ s.t. } \psi^\top \mu(\pi_{\text{exp}}) \geq \psi^\top \mu(\pi) + D(\pi_{\text{exp}}, \pi) \forall \pi$$

Max-Entropy IRL: $P(\tau) \propto \exp(r_\psi(\tau))$. MLE objective – match feature expectations:

$$\nabla_\psi L = \underbrace{\mathbb{E}_{\mathcal{D}_{\text{exp}}}[\nabla_\psi r_\psi]}_{\text{expert features}} - \underbrace{\mathbb{E}_{P(\tau)}[\nabla_\psi r_\psi]}_{\text{model features}}$$

Requires solving RL in inner loop; handles stochastic/suboptimal experts.

GAIL: Adversarial IRL. Jointly optimise:

$$\min_{\pi_\theta} \max_{D_\psi} \mathbb{E}_{\pi_{\text{exp}}}[\log D_\psi(s, a)] + \mathbb{E}_{\pi_\theta}[\log(1 - D_\psi(s, a))]$$

Reward $r(s, a) = -\log(1 - D_\psi(s, a))$; update π_θ via PPO/TRPO.

RLHF:

1. *SFT:* Fine-tune pretrained LM on curated demos.
2. *Collect preferences:* Human feedback.
3. *Reward Model:* $P(y_w \succ y_l) = \sigma(r_\phi(x, y_w) - r_\phi(x, y_l))$. Train via BCE: $\max_\phi \mathbb{E}[\log P(y_w \succ y_l)]$.
4. *RL fine-tune* (PPO) with KL regularisation to prevent reward hacking:
 $\max_{\pi} \mathbb{E}_{(x, y) \sim \pi} [r_\phi(x, y)] - \beta D_{\text{KL}}(\pi \| \pi_{\text{SFT}})$.

RLAIF: Feedback from AI instead of humans.

DPO: Skips explicit reward model. Reparametrise optimal reward via policy ratio:

$r^*(x, y) = \beta \log \frac{\pi^*(y|x)}{\pi_{\text{SFT}}(y|x)} + \beta \log Z(x)$. Substituting into BT loss and dropping $Z(x)$ (cancels):

$$L_{\text{DPO}} = -\mathbb{E} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{SFT}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{SFT}}(y_l|x)} \right) \right]$$

One-stage supervised; no RL loop.

Goodhart’s Law: proxy reward $\neq$ true reward at optimum. Mitigate: KL penalty, ensemble r_ϕ , diverse evals.

15. Planning & Search (Known Model)

MPC (Receding Horizon): Solve $\max_{a_{0:H}} \sum_{t=0}^{H-1} \gamma^t r(s_t, a_t)$; execute a_0^* only; re-plan from new s_1 . Corrects model errors online.

CEM (Cross-Entropy Method): Gradient-free continuous optimiser. Init $\mu_0=0, \Sigma_0=\sigma^2 I$. Each iter: (1) sample N seqs $a_{0:H}^{(i)} \sim \mathcal{N}(\mu_k, \Sigma_k)$; (2) simulate, score; (3) keep top- E elites; (4) refit μ_{k+1}, Σ_{k+1} . Can warm-start from previous μ .

MCTS: 4 phases; builds asymmetric tree toward promising regions.

1. *Selection* (UCT): $a^* = \arg \max_a \left[Q(s, a) + c \sqrt{\frac{\ln N(s)}{N(s, a)}} \right]$.
2. *Expansion:* Add new child node at leaf.
3. *Simulation:* Rollout with fast policy to terminal.
4. *Backpropagate:* $N(s, a) += 1$; $Q(s, a) \leftarrow Q + \frac{1}{N}(G - Q)$.

Play: $a^* = \arg \max_a N(s, a)$ (visit count, not Q).

AlphaGo: MCTS guided by SL+RL policy net p_σ , fast rollout p_π , value net v_θ . *Modified UCT* (PUCT): $a^* = \arg \max_a \left[Q(s, a) + c \frac{p_\sigma(a|s)}{1 + N(s, a)} \right]$.

Leaf evaluation: $V(s_L) = (1 - \lambda)v_\theta(s_L) + \lambda z$, where $z \in \{-1, 0, +1\}$ is rollout outcome.

AlphaZero: Single net $f_\theta(s) = (p, v)$; self-play only; no SL, no rollout policy.

MCTS produces improved policy: $\pi_{\text{MCTS}}(a|s) \propto N(s, a)^{1/\tau}$ ($\tau \rightarrow 0$: greedy; $\tau=1$: proportional). Training targets $(s_t, \pi_{\text{MCTS}}(\cdot|s_t), z)$; loss:

$$L = \underbrace{(v_\theta(s) - z)^2}_{\text{value}} - \underbrace{\pi_{\text{MCTS}}^\top \log p_\theta}_{\text{policy (cross-ent.)}} + \underbrace{c \|\theta\|^2}_{\text{regularization}}$$

16. Model-Based RL (Learned Model)

Learn $\hat{P}_\phi(s_{t+1}|s_t, a_t), \hat{R}_\phi(s, a)$ from interactions. **Trade-off:** Better sample efficiency; but distributional shift + compounding error ($\leq H \cdot \epsilon_\phi$ per step) + model exploitation can hurt asymptotic performance.

Dyna-Q: Interleave real and imagined updates each timestep.

1. Act, observe (s, a, r, s') ; store in $\mathcal{D}$; update model $\hat{P}_\phi$.
2. Real update: $Q \leftarrow Q + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$.

3. Repeat $K \times$: sample $(s, a) \sim \mathcal{D}, (r', \hat{s}') \sim \hat{P}_\phi$, update Q .

$K=0$: Q-learning. Large K : fast convergence if model accurate.

Dyna-Q+: Adds an exploration bonus. $r^+ = r + \kappa \sqrt{\tau}$

MBPO: Ensemble $\hat{P}_\phi$; branch short k -step rollouts ($k=1-5$) from real $s_0 \sim \mathcal{D}_{\text{real}}$; train SAC on mixed buffer $\mathcal{D}_{\text{real}} \cup \mathcal{D}_{\text{model}}$. Limits compounding error; $\sim 20 \times$ sample efficiency.

Probabilistic Ensembles (PETS): M Gaussian models $\hat{P}_{\phi_i}(s'|s, a) = \mathcal{N}(\mu_{\phi_i}, \Sigma_{\phi_i})$.

$$\text{Var}_{\text{ens}} = \underbrace{\frac{1}{M} \sum_i \Sigma_{\phi_i}}_{\text{aleatoric (env. noise)}} + \underbrace{\frac{1}{M} \sum_i (\mu_{\phi_i} - \bar{\mu})^2}_{\text{epistemic (model unc.)}}$$

Use epistemic uncertainty for exploration or pessimistic planning.

Latent World Models: Encode high-dim obs. to compact z_t .

Dreamer: AC (actor π_θ , critic V_ψ) trained entirely in imagination by backpropagating through RSSM dynamics; no real env steps during policy update.

MuZero: No observation decoder. Learns *value-equivalent* model: representation $h_t = e(o_t)$; dynamics $(r, h_{t+1}) = g(h_t, a_t)$; prediction $(p, v) = f(h_t)$. MCTS in latent h space. Only requires accurate value/policy/reward, not s' reconstruction.

17. RL for Foundation Models

LLM as Policy: $\pi_\theta(a_t|s_t) = p_\theta(y_t|x, y_{<t})$. State = prompt + prefix; action = next token/tool call; reward = preference/verifier/task success. Challenges: huge $|A|$ ($\sim 10^5$ tokens), long horizons, sparse seq-level rewards, conflicting objectives, distribution shift.

Pretraining (BC from text): $\max_\theta \sum_{(x, y) \in \mathcal{D}} \sum_t \log \pi_\theta(y_t|x, y_{<t})$. Imitates data distribution; not the same as being a helpful assistant.

SFT: Same MLE objective on curated demos $(x, y^*) \sim \mathcal{D}_{\text{demo}}$. Simple & stable; copies behaviour rather than maximising preference.

RLHF: (1) SFT. (2) Bradley-Terry reward model from human pref. pairs $y_w \succ y_l$: $\mathcal{L}_{\text{RM}} = -\mathbb{E}[\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l))]$. (3) KL-regularised RL:

$$\max_\theta \mathbb{E}_{y \sim \pi_\theta} \left[r_\phi(x, y) - \beta \log \frac{\pi_\theta(y|x)}{\pi_{\text{ref}}(y|x)} \right]$$

KL penalty (π_θ vs π_{ref}) = stay close to SFT globally; PPO clipping (π_θ vs $\pi_{\theta_{\text{old}}}$) = stable local update.

GRPO: No value function. Sample G outputs per prompt, score with r_ϕ , normalise within group: $\hat{A}_i = \frac{r_i - \text{mean}(\mathbf{r}_1; G)}{\text{std}(\mathbf{r}_1; G) + \delta}$. Apply $\hat{A}_i$ to all tokens of y_i as in PPO-clip. Saves memory vs PPO+critic.

RLHF Failure Modes: Reward hacking (proxy $\neq$ intent); overoptimisation (quality drops after too many steps); preference misspec. (style $\succ$ correctness); mode collapse (generic safe answers); RM distribution shift. Mitigate: KL penalty, ensemble r_ϕ , early stopping.

RLVR: Replace learned r_ϕ with automatic checker (math verifier, code tests, game score). Composite reward: $r(y) = \lambda_{\text{fnt}} r_{\text{fnt}}(y) + \lambda_{\text{cor}} r_{\text{cor}}(y)$.

Signal	Strength	Risk
SFT demos	stable/cheap	limited by examples
RLHF human prefs	subj. quality	proxy hacking
DPO pref. pairs	simple/stable	offline limits
RLVR verifier	scalable	narrow/verifier hacking

Generalist Agents: Tokenise text, images, state, actions into shared sequence $\Rightarrow$ LM and action prediction structurally identical. **Gato:** single Transformer for 600+ tasks via supervised seq. modelling on mixed offline data.

18. Taxonomies & Comparisons

Deadly Triad: FA + Bootstrapping + Off-Policy $\Rightarrow$ can diverge.

Aleatoric uncertainty: inherent randomness in env.

Epistemic uncertainty: due to limited knowledge.

	DP	MC	TD
Bootstraps?	Yes	No	Yes
Samples?	No	Yes	Yes
Model req.?	Yes	No	No

	NPG	TRPO	PPO
Policy type	On-policy	On-policy	On-policy
Constraint	Implicit KL	Hard KL	Soft KL
Key idea	Natural grad.	Trust region	Clipping

	PPO	TD3	SAC
Policy	Stochastic	Deterministic	Stochastic
Data	On-policy	Off-policy	Off-policy
Explore	Entropy	Added noise	Entropy obj.

Best Fit	Stable, large-scale	Cont. control baseline	Sample-efficient
Use Cases	Sims, LLMs, games	Benchmarks, precision	Robotics, real-world
Trade-off	Data hungry	Weak exploration	More moving parts

Method	Family	OOD	Stitching
BC Reg	Behavior reg.	Avoided	Limited
CQL	Conservative Q	Penalized	Yes
IQL	Conservative	Avoided	Yes
DT	Seq. modeling	Avoided	No

MCTS		CEM/RS
Actions	Discrete	Continuous
Object	Tree	Sequence
Guidance	UCT+nets	Elite selection

On-policy	REINFORCE, SARSA, PPO
Off-policy	Q-Learning, DQN, DDPG, TD3, SAC
Offline	CQL, IQL, DT, BC

