




**JOHANNES KEPLER
UNIVERSITY LINZ**

Generative and unsupervised DL methods: Generative Adversarial Networks (GANs)



Günter Klambauer
LIT AI Lab & Institute for Machine Learning

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

VAE vs GAN images



VAE



DCGAN

VAE vs GAN images




VAE



DCGAN

Overview of generative, deep models



12.4. Generative adversarial nets (GANs)

- Suggested in 2014 by Ian Goodfellow
- New type of network and loss, in which two networks play against each other
 - Discriminator network: tries to discriminate between real and generated images
 - Generator network: tries to generate realistic images, to trick the discriminator



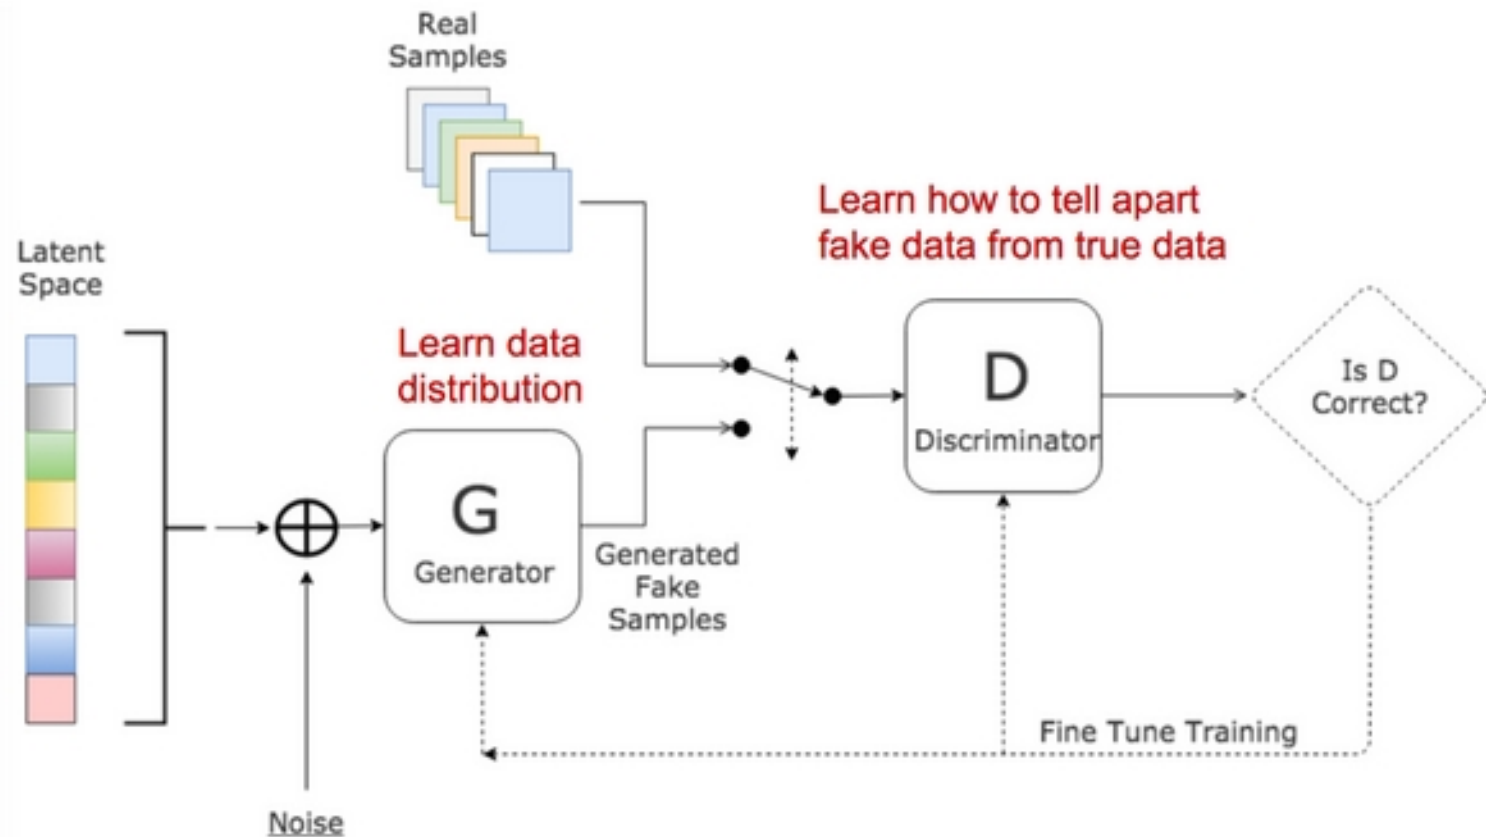
a)




b)

Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. (2014). Generative adversarial nets. In Advances in neural information processing systems (pp. 2672-2680).

The GAN framework



The GAN framework (alternative view)



Objective (loss) function for GANs

$$\begin{aligned}\min_G \max_D \mathcal{L}(D, G) &= \mathbb{E}_{\mathbf{x} \sim p_r} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))] \\ &= \mathbb{E}_{\mathbf{x} \sim p_r} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_g} [\log(1 - D(\mathbf{x}))]\end{aligned}$$

- Two functions have to be optimized
 - Discriminator D
 - Generator G
- Distributions:
 - Random noise $\mathbf{z} \sim p_z$
 - Real world samples / data distribution $\mathbf{x} \sim p_r$
 - Samples from the generator $\mathbf{x} \sim p_g$

Progress of images produced by GANs



Further details on GANs

- Pen & paper lecture

The GAN training algorithm

Algorithm 12.1 GAN training

BEGIN GAN training

while Number of epochs $\leq N$ **do**

 Sample minibatch of size M from data: $\mathbf{x}^1, \dots, \mathbf{x}^M \sim p_r$

 Sample minibatch of size M of noise: $\mathbf{z}^1, \dots, \mathbf{z}^M \sim p_z$

 Take a gradient descent step on the generator parameters θ :

$$\nabla_{\theta} \mathcal{L}(D_w, G_{\theta}) = \frac{1}{M} \nabla_{\theta} \sum_{i=1}^M \log(1 - G_{\theta}(D_w(\mathbf{z}^i)))$$

 Take a gradient ascent step on the discriminator parameters w :

$$\nabla_w \mathcal{L}(D_w, G_{\theta}) = \frac{1}{M} \nabla_w \sum_{i=1}^M \log(D_w(\mathbf{x}^i)) + \log(1 - G_{\theta}(D_w(\mathbf{z}^i)))$$

end while

END GAN training
