




**JOHANNES KEPLER
UNIVERSITY LINZ**

Bayesian Deep Learning & Uncertainty estimation



Günter Klambauer
LIT AI Lab & Institute for Machine Learning
@gklambauer



This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

Bayesian Deep Learning – why?

- Ability to obtain uncertainty estimates
 - Improved decision making
 - Increased trust in ML systems
 - Responsible deployment of ML systems
 - Detecting unreliable inputs

Notation

- The machine learning model $g(\mathbf{x}; \mathbf{w})$ will be $p(\mathbf{y} \mid \mathbf{w}, \mathbf{x})$ or at times $p_{\mathbf{w}}(\mathbf{y} \mid \mathbf{x})$.
- We will use $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_N, \mathbf{y}_N)\}$ to denote the labeled training data. In other lectures, the notation $\mathbf{z} = (\mathbf{x}, \mathbf{y})$ for a pair and $\mathbf{Z} = (\mathbf{X}, \mathbf{y})$ for the full supervised dataset is used – we switch to $\mathcal{D}$ to avoid confusion with latent variables $\mathbf{z}$ or $\mathbf{Z}$
- The prediction $\hat{y}$ will become $\hat{p}$ or $\hat{p}_k$.
- A new data point that we aim to predict will be denoted as $(\mathbf{x}^{\text{test}}, \mathbf{y}^{\text{test}})$.
- A model will be denoted with $\mathcal{M}$

Math background

- Basic knowledge about Bayes Theorem and its use for parameter estimation for ML models; Bayesian model average
- Definitions of
 - Entropy H
 - Cross-Entropy CE
 - Mutual information I
 - KL divergence KL

Likelihood, Prior, Posterior, Evidence

- **Maximum likelihood:** find the most likely parameters given data

$$L(\mathcal{D}; \mathbf{w}) = p(\mathcal{D} \mid \mathbf{w})$$

- However, for complex models this leads to overfitting. Therefore, some $\mathbf{w}$ should be more probable than others

- Prior distribution over parameters: $p(\mathbf{w})$

- Bayes formula $p(\mathbf{w} \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \mathbf{w}) p(\mathbf{w})}{p_{\mathbf{w}}(\mathcal{D})}$

- Important note! Usually: $p(\mathcal{D}) \neq p_{\mathbf{w}}(\mathcal{D})$

- Maximum A Posteriori: (last term does not depend on params)

$$-\log p(\mathbf{w} \mid \mathcal{D}) = -\log p(\mathcal{D} \mid \mathbf{w}) - \log p(\mathbf{w}) + \log p_{\mathbf{w}}(\mathcal{D})$$

Maximum A Posteriori (MAP) approach for Deep Nets: a usual case

- Usually, the aim is to find the most likely parameters given data

$$\begin{aligned}\tilde{\boldsymbol{w}} &= \operatorname{argmax}_{\boldsymbol{w}} p(\boldsymbol{w} \mid \boldsymbol{x}, \boldsymbol{y}) \\ &= \operatorname{argmin}_{\boldsymbol{w}} -\log p(\boldsymbol{y} \mid \boldsymbol{x}, \boldsymbol{w}) - \log p(\boldsymbol{w})\end{aligned}$$

- For example, softmax & cross-entropy at output and L2 regularization lead to

$$\tilde{\boldsymbol{w}} = \operatorname{argmin}_{\boldsymbol{w}} \sum_k y_k \log(\hat{p}_k) + \lambda \|\boldsymbol{w}\|^2$$

- Note: $p(\boldsymbol{y} \mid \boldsymbol{x}, \boldsymbol{w})$ represents uncertainty about the label. However, we can make only a single prediction with the parameters we found
- Still a point estimate! No full Bayesian treatment! We want a full distribution over parameters!

A Bayesian approach: a distribution of parameters is wanted; supervised!

- Bayesian setting:

$$p(\mathbf{w} \mid \mathbf{x}, \mathbf{y}) = \frac{p(\mathbf{y} \mid \mathbf{w}, \mathbf{x}) \cdot p(\mathbf{w})}{\int_{\mathcal{W}} p(\mathbf{y} \mid \mathbf{w}, \mathbf{x}) \cdot p(\mathbf{w}) d\mathbf{w}} = \frac{p(\mathbf{y} \mid \mathbf{w}, \mathbf{x}) \cdot p(\mathbf{w})}{p(\mathbf{y} \mid \mathbf{x})}$$

- Posterior: $p(\mathbf{w} \mid \mathbf{x}, \mathbf{y})$
- Prior: $p(\mathbf{w})$
- Likelihood: $p(\mathbf{y} \mid \mathbf{w}, \mathbf{x})$
- Evidence: $\int_{\mathcal{W}} p(\mathbf{y} \mid \mathbf{w}, \mathbf{x}) \cdot p(\mathbf{w}) d\mathbf{w}$
 - Usually intractable; especially for DNNs
 - If all other distributions are Gaussian: tractable. “Gaussian process”

Example: Gaussian data, Gaussian prior

- Data: $x \mid \mu \sim \mathcal{N}(\mu, \sigma_x^2)$ Pior: $\mu \sim \mathcal{N}(\nu, \sigma_\mu^2)$
- Parameter posterior

$$p(\mu \mid x) = \frac{p(x \mid \mu)p(\mu)}{\int_{-\infty}^{\infty} p(x \mid m)p(m)dm} \propto p(x \mid \mu)p(\mu)$$

$$p(x \mid \mu)p(\mu) = \frac{1}{\sqrt{2\pi\sigma_x^2}} e^{-\frac{(x-\mu)^2}{2\sigma_x^2}} \frac{1}{\sqrt{2\pi\sigma_\mu^2}} e^{-\frac{(\mu-\nu)^2}{2\sigma_\mu^2}}$$

- **Completing squares in exponent** and assuming normalized distributions, we see that this is again a Gaussian

$$\mu \mid x \sim \mathcal{N} \left(\frac{\mu\sigma_\mu^2 + \nu\sigma_x^2}{\sigma_\mu^2 + \sigma_x^2}, \frac{1}{1/\sigma_x^2 + 1/\sigma_\mu^2} \right)$$

Intuition?


- Assume you train a vision model $p(y \mid x, w)$ that should predict binary classes (“cat” vs “dog”).
- It provides for a particular image:

$$p(y = 1 \mid x, w) = 0.5$$

- Is the model uncertain about its prediction?
- Correct answer: we don’t know!
 - **A)** It could be the perfect model, but the image is ambiguous the prediction is correct: every time one labels the image it is labeled with 50% cat and 50% dog; → **aleatoric uncertainty**
 - **B)** The model is garbage and the image is indeed shows a cat (or: dog); uncertainty about the parameters → **epistemic uncertainty**
- **Does the model not know anything (B)? Or**
- **Does it know that it does not know anything (A).**

Types of uncertainty


- Levels of uncertainty representations




- Note: the softmax output of a neural network “pretends” to have no *epistemic uncertainty* and to just provide *aleatoric uncertainty*

Types of uncertainty


- Levels of uncertainty representations



(a) Confident Prediction



(b) High data uncertainty




(c) Out-of-distribution

- Note: the softmax output of a neural network “pretends” to have no *epistemic uncertainty* and to just provide *aleatoric uncertainty*

Sources of uncertainty


- Goal: have a **good model** for $p(\mathbf{y} | \mathbf{x})$; this would give us **aleatoric uncertainty**
 - Model $p(\mathbf{y} | \mathbf{x}; \mathbf{w})$
 - Model will never be perfect → **epistemic uncertainty** wanted
- Sources of epistemic uncertainty
 - A) No access to $p(\mathbf{x}, \mathbf{y})$, thus to get $\mathbf{w}$, we need data
 - Drawing data from joint distribution $D \simeq p(\mathbf{x}, \mathbf{y})$
 - Uncertainty in the drawn data (usually: no control over; IID)
 - B) Uncertainty about definition of hypothesis space
 - Usually hard to model
 - Expressive enough? Often has to be assumed
 - Assumption: $p(\mathbf{y} | \mathbf{x}) = p(\mathbf{y} | \mathbf{x}; \mathbf{w}^*)$; $\mathbf{w}^*$ is the parameter-vector of the true model
 - C) Uncertainty about chosen parameters $\mathbf{w}$
 - Bayesian framework: posterior distribution
 - Consider all possible $\mathbf{w}$, weighted by posterior $p(\mathbf{w} | D)$ → intractable; Methods needed

B) Neural network generalization




- How is model class performance ($\sim$ inductive bias) distributed over the range of all possible datasets (support)


B)→ C) A probabilistic perspective of generalization



(b)



(c)



(d)

Key difference between frequentist and Bayesian approach: marginalization

- Posterior predictive distribution / distribution of outputs:

$$\begin{aligned} p(\mathbf{y}^{\text{test}} \mid \mathbf{x}^{\text{test}}, \mathcal{D}) &= \mathbb{E}_{\mathbf{w} \sim p(\mathbf{w} \mid \mathcal{D})} [p(\mathbf{y}^{\text{test}} \mid \mathbf{w}, \mathbf{x}^{\text{test}})] , \\ &= \int_{\mathcal{W}} p(\mathbf{y}^{\text{test}} \mid \mathbf{w}, \mathbf{x}^{\text{test}}) p(\mathbf{w} \mid \mathcal{D}) d\mathbf{w} \end{aligned}$$

- Defines probability for class label given input and dataset
- Marginalization over parameters
- “Bayesian model averaging”

Inference methods

- Maximum likelihood

$$\boldsymbol{w}_{\text{ML}} = \operatorname{argmax}_{\boldsymbol{w}} p(\mathcal{D} \mid \boldsymbol{w})$$

- Maximum A-posteriori (MAP) inference

$$\boldsymbol{w}_{\text{MAP}} = \operatorname{argmax}_{\boldsymbol{w}} p(\boldsymbol{w} \mid \mathcal{D}) = \operatorname{argmax}_{\boldsymbol{w}} p(\mathcal{D} \mid \boldsymbol{w})p(\boldsymbol{w})$$

- Bayesian and approximate Bayesian inference

$$\int_{\boldsymbol{w}} F(\boldsymbol{w})p(\boldsymbol{w} \mid \mathcal{D})d\boldsymbol{w} \approx \frac{1}{M} \sum_{m=1}^M F(\boldsymbol{w}_k); \quad \boldsymbol{w}_k \sim p(\boldsymbol{w} \mid \mathcal{D})$$

Definition of uncertainty

- *Total uncertainty* often defined as entropy of the predictive distribution (e.g. Gal, 2016; Hüllermeier, 2021):

$$H[p(\mathbf{y} \mid \mathbf{x}, \mathcal{D})] = \int_{\mathcal{W}} H[p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})] p(\tilde{\mathbf{w}} \mid \mathcal{D}) d\tilde{\mathbf{w}} + I(Y, W) .$$

- The *aleatoric uncertainty* is given by

$$\mathbb{E}_{p(\tilde{\mathbf{w}} \mid \mathcal{D})} [H[p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})]] = \int_{\mathcal{W}} H[p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})] p(\tilde{\mathbf{w}} \mid \mathcal{D}) d\tilde{\mathbf{w}}$$

- *Epistemic uncertainty* obtained as difference

$$\begin{aligned} I(Y; W) &= H[p(\mathbf{y} \mid \mathbf{x}, \mathcal{D})] - \mathbb{E}_{p(\tilde{\mathbf{w}} \mid \mathcal{D})} [H[p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})]] \\ &= \int_{\mathcal{W}} (H[p(\mathbf{y} \mid \mathbf{x}, \mathcal{D})] - H[p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})]) p(\tilde{\mathbf{w}} \mid \mathcal{D}) d\tilde{\mathbf{w}} \end{aligned}$$

$$\text{JKU JOHANNES KEPLER UNIVERSITY LINZ} \quad p(\mathbf{y} \mid \mathbf{x}, \mathcal{D}) = p(\mathbf{y} \mid \mathbf{x}, \mathbf{w}^*) = \int_{\mathcal{W}} p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}}) p(\tilde{\mathbf{w}} \mid \mathcal{D}) d\tilde{\mathbf{w}}$$

Intuition?

- Assume you train a vision model $p(y \mid x, w)$ that should predict binary classes (“cat” vs “dog”).
- It provides for a particular image:

$$p(y = 1 \mid x, w) = 0.5$$

- Is the model uncertain about it’s prediction?
- Correct answer: we don’t know!
 - A) It could be the perfect model, but the image is ambiguous the prediction is correct: every time one labels the image it is labeled with 50% cat and 50% dog; → **aleatoric uncertainty**
 - B) The model is garbage and the image is indeed shows a cat (or: dog); uncertainty about the parameters → **epistemic uncertainty**

Monte-Carlo Approximation


- How to approximate the integral? → sampling from posterior

$$p(\mathbf{y} \mid \mathbf{x}, \mathcal{D}) = p(\mathbf{y} \mid \mathbf{x}, \mathbf{w}^*) = \int_{\mathcal{W}} p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}}) p(\tilde{\mathbf{w}} \mid \mathcal{D}) \, d\tilde{\mathbf{w}} \approx \frac{1}{K} \sum_{k=1}^K p(\mathbf{y} \mid \mathbf{x}, \mathbf{w}_k); \quad \mathbf{w}_k \sim p(\tilde{\mathbf{w}} \mid \mathcal{D})$$

- How to sample from posterior?
 - METHODS (see later)

Full predictive distribution

- Linear regression example
- Uncertainty close to data points is small
- High uncertainty elsewhere



A selection of Bayesian DL approaches

- 15.1. Deep Ensembles
- 15.2. Weight uncertainty in DNNs (variational approach)
- 15.3. Monte-Carlo Dropout
- 15.4. Hessian-based approach

15.1. Deep Ensembles

- Main idea: approximate expectation with averages

$$p(\mathbf{y}^{\text{test}} \mid \mathbf{x}^{\text{test}}, \mathcal{D}) = \mathbb{E}_{\mathbf{w} \sim p(\mathbf{w} \mid \mathcal{D})} [p(\mathbf{y}^{\text{test}} \mid \mathbf{w}, \mathbf{x}^{\text{test}})]$$

- Train several DNNs with different random seeds, obtain parameters $\mathbf{w}^{(m)}$ and take average

$$p(\mathbf{y}^{\text{test}} \mid \mathbf{x}^{\text{test}}, \mathcal{D}) \approx \frac{1}{M} \sum_{m=1}^M p(\mathbf{y}^{\text{test}} \mid \mathbf{w}^{(m)}, \mathbf{x}^{\text{test}},)$$

- Averages can be weighted differently
- Appears as simple technique, but performs well in practice

Deep Ensembles

Algorithm 1 Pseudocode of the training procedure for our method

- 1: $\triangleright$ Let each neural network parametrize a distribution over the outputs, i.e. $p_{\theta}(y|\mathbf{x})$. Use a proper scoring rule as the training criterion $\ell(\theta, \mathbf{x}, y)$. Recommended default values are $M = 5$ and $\epsilon = 1\%$ of the input range of the corresponding dimension (e.g 2.55 if input range is $[0, 255]$).
 - 2: Initialize $\theta_1, \theta_2, \dots, \theta_M$ randomly
 - 3: **for** $m = 1 : M$ **do** $\triangleright$ train networks independently in parallel
 - 4: Sample data point n_m randomly for each net $\triangleright$ single n_m for clarity, minibatch in practice
 - 5: Generate adversarial example using $\mathbf{x}'_{n_m} = \mathbf{x}_{n_m} + \epsilon \text{sign}(\nabla_{\mathbf{x}_{n_m}} \ell(\theta_m, \mathbf{x}_{n_m}, y_{n_m}))$
 - 6: Minimize $\ell(\theta_m, \mathbf{x}_{n_m}, y_{n_m}) + \ell(\theta_m, \mathbf{x}'_{n_m}, y_{n_m})$ w.r.t. θ_m $\triangleright$ adversarial training (optional)
-


Notation in code
does not
fit to our notation!

15.2. Variational approach: Weight Uncertainty in DNNs

- Approximate predictive distribution
- Approach: Weight uncertainty in neural networks (Blundell et al., 2015).
- Variational approach, learning finds the parameters w of a parameterized distribution of the weights $q(w | \theta)$
- New new set of parameters θ that we did not have before.
 - parameters determine the distribution of weights,
 - e.g. mean and standard deviation of Gaussian

Weight Uncertainty in DNNs

Intuition



Weight Uncertainty in DNNs

- Gaussian approximation of the parameter posterior
- Minimization of KL divergence with true posterior

$$\begin{aligned}\tilde{\boldsymbol{\theta}} &= \operatorname{argmin}_{\boldsymbol{\theta}} \text{KL}(q(\boldsymbol{w} \mid \boldsymbol{\theta}) \parallel p(\boldsymbol{w} \mid \mathcal{D})) \\ &= \operatorname{argmin}_{\boldsymbol{\theta}} \int_{\boldsymbol{w}} q(\boldsymbol{w} \mid \boldsymbol{\theta}) \log \frac{q(\boldsymbol{w} \mid \boldsymbol{\theta})}{p(\boldsymbol{w})p(\mathcal{D} \mid \boldsymbol{w})} d\boldsymbol{w} \\ &= \operatorname{argmin}_{\boldsymbol{\theta}} \text{KL}(q(\boldsymbol{w} \mid \boldsymbol{\theta}) \parallel p(\boldsymbol{w})) - \mathbb{E}_{\boldsymbol{w} \sim q(\boldsymbol{w} \mid \boldsymbol{\theta})} [\log p(\mathcal{D} \mid \boldsymbol{w})]\end{aligned}$$

Weight Uncertainty in DNNs

- Objective function (ELBO):

$$\begin{aligned} L(\boldsymbol{\theta}, \mathcal{D}) &= L(\boldsymbol{\theta}, \mathbf{X}, \mathbf{Y}) = \text{KL} (q(\mathbf{w} \mid \boldsymbol{\theta}) \parallel p(\mathbf{w})) - \mathbb{E}_{\mathbf{w} \sim q(\mathbf{w} \mid \boldsymbol{\theta})} [\log p(\mathcal{D} \mid \mathbf{w})] \\ &= \mathbb{E}_{\mathbf{w} \sim q(\mathbf{w} \mid \boldsymbol{\theta})} [\log q(\mathbf{w} \mid \boldsymbol{\theta}) - \log p(\mathbf{w})] - \mathbb{E}_{\mathbf{w} \sim q(\mathbf{w} \mid \boldsymbol{\theta})} [\log p(\mathcal{D} \mid \mathbf{w})] \end{aligned}$$

- Similar to objective in VAEs
- Involves sampling from posterior estimates
- How to approach this with gradient descent? We cannot backpropagate to a sampling procedure

Weight Uncertainty in DNNs: Reparametrization trick

- We assume that weights are determined by some function that has both a deterministic and a random part:

$$w = h(\theta, \epsilon)$$

- ϵ : noise
- θ : parameters of probability distribution
- Similar to reparametrization trick in VAEs
- E.g. mean and variance of Gaussian plus standard noise

Weight Uncertainty in DNNs: Reparametrization trick

- Backpropagation to parameters of probability distribution

$$\begin{aligned}\frac{\partial}{\partial \boldsymbol{\theta}} \mathbb{E}_{\boldsymbol{w} \sim q(\boldsymbol{w} | \boldsymbol{\theta})} [f(\boldsymbol{w}, \boldsymbol{\theta})] &= \frac{\partial}{\partial \boldsymbol{\theta}} \int f(\boldsymbol{w}, \boldsymbol{\theta}) q(\boldsymbol{w} | \boldsymbol{\theta}) d\boldsymbol{w} = \\ &= \frac{\partial}{\partial \boldsymbol{\theta}} \int f(\boldsymbol{w}, \boldsymbol{\theta}) q(\boldsymbol{\epsilon}) d\boldsymbol{\epsilon} \\ &= \mathbb{E}_{\boldsymbol{\epsilon} \sim q(\boldsymbol{\epsilon})} \left[\frac{\partial f(\boldsymbol{w}, \boldsymbol{\theta})}{\partial \boldsymbol{w}} \frac{\partial \boldsymbol{w}}{\partial \boldsymbol{\theta}} + \frac{\partial f(\boldsymbol{w}, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \right]\end{aligned}$$

- We used Leibnitz integral rule to switch integration and differentiation
- We use: $f(\boldsymbol{w}, \boldsymbol{\theta}) = \log q(\boldsymbol{w} | \boldsymbol{\theta}) - \log p(\boldsymbol{w}) - \log p(\mathcal{D} | \boldsymbol{w})$

Weight Uncertainty in DNNs: Reparametrization trick

- With $f(\mathbf{w}, \boldsymbol{\theta}) = \log q(\mathbf{w} \mid \boldsymbol{\theta}) - \log p(\mathbf{w}) - \log p(\mathcal{D} \mid \mathbf{w})$ we approximate the objective function by

$$L(\boldsymbol{\theta}, \mathcal{D}) \approx \sum_{m=1}^M \log q(\mathbf{w}^{(m)} \mid \boldsymbol{\theta}) - \log p(\mathbf{w}^{(m)}) - \underbrace{\log p(\mathcal{D} \mid \mathbf{w}^{(m)})}_{\text{usual obj of a DNN}}$$

- Where $\mathbf{w}^{(m)}$ is a parameter vector randomly drawn from the estimated posterior distribution $q(\mathbf{w} \mid \boldsymbol{\theta})$

Weight Uncertainty in DNNs: Algorithm

- Sample $\varepsilon \sim \mathcal{N}(0, 1)$,
- set $\mathbf{w} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \varepsilon$
- let $\boldsymbol{\theta} = (\boldsymbol{\mu}, \boldsymbol{\sigma})$,
- let $f(\mathbf{w}, \boldsymbol{\theta}) = \log q(\mathbf{w} \mid \boldsymbol{\theta}) - \log p(\mathbf{w}) - \log p(\mathcal{D} \mid \mathbf{w})$,
- calculate gradients of $f(\mathbf{w}, \boldsymbol{\theta})$ with respect to $\boldsymbol{\mu}$ using backprop:

$$\Delta_{\boldsymbol{\mu}} = \frac{\partial f(\mathbf{w}, \boldsymbol{\theta})}{\partial \mathbf{w}} \odot \mathbf{1} + \frac{\partial f(\mathbf{w}, \boldsymbol{\theta})}{\partial \boldsymbol{\mu}}$$

- calculate gradients of $f(\mathbf{w}, \boldsymbol{\theta})$ with respect to $\boldsymbol{\sigma}$ using backprop:


$$\Delta_{\boldsymbol{\sigma}} = \frac{\partial f(\mathbf{w}, \boldsymbol{\theta})}{\partial \mathbf{w}} \odot \varepsilon + \frac{\partial f(\mathbf{w}, \boldsymbol{\theta})}{\partial \boldsymbol{\sigma}}$$

- and update the variational parameters:


$$\begin{aligned}\boldsymbol{\mu}^{\text{new}} &= \boldsymbol{\mu}^{\text{old}} - \eta \Delta_{\boldsymbol{\mu}} \\ \boldsymbol{\sigma}^{\text{new}} &= \boldsymbol{\sigma}^{\text{old}} - \eta \Delta_{\boldsymbol{\sigma}}\end{aligned}$$

These are
the usual gradients
from backprop

15.3. Monte Carlo Dropout



(a) Standard Neural Net



(b) After applying dropout.

Monte-Carlo Dropout

- Similar derivation to *Weight Uncertainty in DNNs*
- Use of the equivalence to Gaussian processes
- Consider a two-layer network
- Intractable posterior; approximation with $q(\mathbf{w})$
 - Weights assumed to come from mixture

$$q(\mathbf{W}_1) = \prod_{q=1}^Q q(\mathbf{w}_q),$$

$$q(\mathbf{w}_q) = p_1 \mathcal{N}(\mathbf{m}_q, \sigma^2 \mathbf{I}_K) + (1 - p_1) \mathcal{N}(0, \sigma^2 \mathbf{I}_K)$$

- Applying Dropout is like drawing from this posterior estimates
- Objective is again ELBO approximated with sampling

MC Dropout

- Two-layer network (one hidden layer):

$$p(y \mid \mathbf{x}, \mathbf{W}^{[1]}, \mathbf{W}^{[2]}) = \text{softmax}(\mathbf{W}^{[2]} f(\mathbf{W}^{[1]} \mathbf{x}))$$

- Gaussian prior on weights

$$p(\mathbf{W}^{[1]}) = \mathcal{N}(0, \mathbf{I}) \quad \text{and} \quad p(\mathbf{W}^{[2]}) = \mathcal{N}(0, \mathbf{I})$$

- Full predictive distribution

$$p(\mathbf{y}^{\text{test}} \mid \mathbf{x}^{\text{test}}, \mathcal{D}) = \int_{\mathcal{W}} p(\mathbf{y}^{\text{test}} \mid \mathbf{W}^{[1]}, \mathbf{W}^{[2]}, \mathbf{x}^{\text{test}}) \underbrace{p(\mathbf{W}^{[1]}, \mathbf{W}^{[2]} \mid \mathcal{D})}_{\text{untractable}} d\mathbf{W}^{[1]} d\mathbf{W}^{[2]}$$

- Minimizing KL-divergence and connection to ELBO (see VAEs)

$$D_{KL}(q(\mathbf{W}^{[1]}, \mathbf{W}^{[2]}) \parallel p(\mathbf{W}^{[1]}, \mathbf{W}^{[2]} \mid \mathcal{D}))$$

$$D_{KL}(q(\mathbf{W}^{[1]}, \mathbf{W}^{[2]}) \parallel p(\mathbf{W}^{[1]})p(\mathbf{W}^{[2]})) - \sum_{n=1}^N \int q(\mathbf{W}^{[1]}, \mathbf{W}^{[2]}) \log(y_n \mid x_n, \mathbf{W}^{[1]}, \mathbf{W}^{[2]})$$

- With variational distribution; mixture of Gaussians

$$q(\mathbf{W}^{[1]}, \mathbf{W}^{[2]}) = q(\mathbf{W}^{[1]})q(\mathbf{W}^{[2]})$$

$$q(\mathbf{W}) = p\mathcal{N}(\mathbf{M}, \sigma^2 \mathbf{I}) + (1 - p)\mathcal{N}(0, \sigma^2 \mathbf{I})$$

MC Dropout

- KL divergence between variational and prior

$$D_{KL}(q(\mathbf{W}^{[l]})||p(\mathbf{W}^{[l]})) \approx dI(\sigma^2 - \log(\sigma^2) - 1) + p/2||\mathbf{W}^{[l]}||_2^2 + \text{const}$$

- Hence: weight matrices sampled from

$$\mathbf{W}^{[l]} \approx \text{diag}(\mathbf{z}^{[l]})\mathbf{M}^{[l]}, \mathbf{z} \sim \text{Bernoulli}(p)$$

- Objective:


$$D_{KL}(q(\mathbf{W}^{[1]}, \mathbf{W}^{[2]}||p(\mathbf{W}^{[1]}, \mathbf{W}^{[2]} | \mathcal{D})) \leq - \sum_{n=1}^N \log p(y_n | x_n, \mathbf{W}_n^{[1]}, \mathbf{W}_n^{[2]}) + p^{[1]}/2||\mathbf{M}_1||_2^2 + p^{[2]}/2||\mathbf{M}_2||_2^2$$

where $\mathbf{W}_n = \text{diag}(\mathbf{z}_n)\mathbf{M}$, $\mathbf{z} \sim \text{Bernoulli}(p)$ are Monte Carlo samples


Same loss as used when training with dropout! (and reg.)

Bayesian DL methods compared


MC Dropout MFVI Deep Ensemble Deterministic Ensemble MC Dropout random referral




(a) test data, AUC



(b) test data, accuracy



(c) distribution shift, AUC



(d) distribution shift, accuracy

Filos, A., Farquhar, S., Gomez, A. N., Rudner, T. G., Kenton, Z., Smith, L., ... & Gal, Y. (2019). A systematic comparison of bayesian deep learning robustness in diabetic retinopathy tasks. arXiv preprint arXiv:1912.10481.

15.4. Hessian-based approach

- Regression task: Gaussian noise on labels

$$p(y \mid \mathbf{x}, \mathbf{w}, \sigma_\epsilon^2) = \mathcal{N}(y \mid g(\mathbf{x}; \mathbf{w}), \sigma_\epsilon^2)$$

- Gaussian prior on weights

$$p(\mathbf{w} \mid \sigma_w^2) = \mathcal{N}(\mathbf{w} \mid \mathbf{0}, \sigma_w^2)$$

- Dataset $\mathcal{D}$ of data points and labels; posterior and log post.:

$$p(\mathbf{w} \mid \mathcal{D}, \sigma_w^2, \sigma_\epsilon^2) \propto p(\mathbf{w} \mid \sigma_w^2) p(\mathcal{D} \mid \mathbf{w}, \sigma_\epsilon^2)$$

$$\log p(\mathbf{w} \mid \mathcal{D}, \sigma_w^2, \sigma_\epsilon^2) = \frac{1}{2\sigma_w^2} \mathbf{w}^T \mathbf{w} - \frac{1}{2\sigma_\epsilon^2} \sum_{n=1}^N (g(\mathbf{x}; \mathbf{w}) - y_n)^2 + C$$

The last function is optimized with the usual gradient descent techniques and we obtain parameters: $\mathbf{w}_{\text{MAP}}$

15.4. Hessian-based approach

- Having found $\mathbf{w}_{\text{MAP}}$, we can make a local Gaussian approximation by evaluating the matrix of second derivatives of the negative log posterior:

$$\mathbf{A} = -\nabla^2 \log p(\mathbf{w} \mid \mathcal{D}, \sigma_w^2, \sigma_\epsilon^2) = 1/\sigma_w^2 \mathbf{I} + 1/\sigma_\epsilon^2 \mathbf{H}$$

where $\mathbf{H}$ is the usual Hessian of the loss function of the neural network. Can be approximated in various ways for small networks.

- Gaussian approximation of posterior

$$q(\mathbf{w} \mid \mathcal{D}) = \mathcal{N}(\mathbf{w} \mid \mathbf{w}_{\text{MAP}}, \mathbf{A}^{-1})$$

Intractable because distribution on weights is inside the neural net

- Predictive distribution (still intractable, sampling from approx)

$$p(y \mid \mathbf{x}, \mathcal{D}) = \int \mathcal{N}(y \mid g(\mathbf{x}, \mathbf{w}), \sigma_\epsilon^2) q(\mathbf{w} \mid \mathcal{D}) d\mathbf{w}$$

15.5. Hessian-based approach

- If we would have a linear model instead of $g(\mathbf{x}; \mathbf{w})$, this integral would be analytically tractable.
- Let us approximate the neural net around $\mathbf{w}_{\text{MAP}}$ with a Taylor approximation up to first degree

$$g(\mathbf{x}; \mathbf{w}) \approx g(\mathbf{x}; \mathbf{w}_{\text{MAP}}) + \mathbf{g}^T (\mathbf{w} - \mathbf{w}_{\text{MAP}})$$

where $\mathbf{g} = \nabla_{\mathbf{w}} g(\mathbf{x}; \mathbf{w})|_{\mathbf{w}_{\text{MAP}}}$

- Linear Gaussian model with Gaussian distribution of weights and prior whose mean is linear function of weights


$$p(y \mid \mathbf{x}, \mathbf{w}, \sigma_{\epsilon}^2) \approx \mathcal{N}(y \mid g(\mathbf{x}, \mathbf{w}_{\text{MAP}}) + \mathbf{g}^T (\mathbf{w} - \mathbf{w}_{\text{MAP}}), \sigma_{\epsilon}^2)$$

General results on marginal Gaussian provide:

$$p(y \mid \mathbf{x}, \mathcal{D}, \sigma_{\epsilon}^2, \sigma_w^2) \approx \mathcal{N}(y \mid g(\mathbf{x}, \mathbf{w}_{\text{MAP}}), \sigma_{\epsilon}^2 + \mathbf{g}^T \mathbf{A} \mathbf{g})$$


Methods developed at JKU: Quantification of Uncertainty with Adversarial Models

- **Problem:** how to find diverse models
 - Methods usually find similar models (through sgd)



- **Idea:** explicitly search for alternative models that also explain the data well (high posterior)
- Searching for modes in posterior
 - **Adversarial attacks** in weight space

Methods developed at JKU: Quantification of Uncertainty with Adversarial Models



Methods developed at JKU:

Quantification of Uncertainty with Adversarial Models

- Definition of *adversarial model*:


Definition 1. Given are a new test data point $\mathbf{x}$, a reference conditional probability model $p(\mathbf{y} \mid \mathbf{x}, \mathbf{w})$ from a model class parameterized by $\mathbf{w}$, a divergence or distance measure $D(\cdot, \cdot)$ for probability distributions, $\gamma > 0$, $\Lambda > 0$, and a dataset $\mathcal{D}$. Then a model with parameters $\tilde{\mathbf{w}}$ that satisfies the inequalities $|\log p(\mathbf{w} \mid \mathcal{D}) - \log p(\tilde{\mathbf{w}} \mid \mathcal{D})| \leq \gamma$ and $D(p(\mathbf{y} \mid \mathbf{x}, \mathbf{w}), p(\mathbf{y} \mid \mathbf{x}, \tilde{\mathbf{w}})) \geq \Lambda$ is called an (γ, Λ) -adversarial model.

- Search for adversarial model:


$$\max_{\delta \in \Delta} D(p(\mathbf{y} \mid \mathbf{x}, \mathbf{w}), p(\mathbf{y} \mid \mathbf{x}, \mathbf{w} + \delta)) \quad \text{s.t.} \quad \log p(\mathbf{w} \mid \mathcal{D}) - \log p(\mathbf{w} + \delta \mid \mathcal{D}) \leq \gamma.$$

- Iterative procedure (gradient-based) employed to search for these models


Methods developed at JKU: Quantification of Uncertainty with Adversarial Models




(a) **Ground Truth - HMC**




(b) **cSG-HMC**




(c) **Laplace**



(d) **MC dropout**



(e) **Deep Ensembles**



(f) **Our Method - QUAM**

Misclassification/OOD detection: different Bayesian DL methods

Table 2: ImageNet-1K results: AUROC using the epistemic uncertainty of a given, pre-selected model (as in Eq. (2)) to distinguish between ID (ImageNet-1K) and OOD samples. Furthermore, we report the AUROC when using the epistemic uncertainty for misclassification detection and the AUC of accuracy over fraction of retained predictions on the ImageNet-1K validation dataset. We also report results for all experiments, using the aleatoric uncertainty of the reference model (Reference).

$\mathcal{D}_{\text{ood}}$ // Task	Reference	cSG-HMC	MCD	DE (LL)	DE (all)	QUAM
ImageNet-O	.626 $\pm$.004	.677 $\pm$.005	.680 $\pm$.003	.562 $\pm$.004	.709 $\pm$.005	.753 $\pm$.011
ImageNet-A	.792 $\pm$.002	.799 $\pm$.001	.827 $\pm$.002	.686 $\pm$.001	.874 $\pm$.004	.872 $\pm$.003
Misclassification	.867 $\pm$.007	.772 $\pm$.011	.796 $\pm$.014	.657 $\pm$.009	.780 $\pm$.009	.904 $\pm$.008
Selective prediction	.958 $\pm$.003	.931 $\pm$.003	.935 $\pm$.006	.911 $\pm$.004	.950 $\pm$.002	.969 $\pm$.002

Summary

- We have investigated Bayesian Deep Learning approaches
- Discussed main difference to frequentist approaches
- Can capture uncertainty about parameters
- Usually approximation of parameter posterior
- Variational approach, Monte-Carlo Dropout, Deep Ensembles

Infinite Width Bayesian Deep Networks are Gaussian Processes

- We revisit linear regression:

$$g(\mathbf{x}) = \hat{y} = \mathbf{w}^T f(\mathbf{x}) = \mathbf{w}^T \mathbf{h}$$

- Now we assume a Gaussian prior over the parameters

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w} \mid \mathbf{0}, \sigma^2 \mathbf{I})$$

- We stack representations and labels of different samples

$$\hat{\mathbf{y}} = \mathbf{H} \mathbf{w}$$

- Then we calculate mean and covariance to define $\mathbf{K}$

$$\mathbb{E}[\hat{\mathbf{y}}] = \mathbf{H} \mathbb{E}[\mathbf{w}] = \mathbf{0}$$

$$\mathbf{J} \text{ Cov}[\hat{\mathbf{y}}] = \mathbb{E}[\hat{\mathbf{y}} \hat{\mathbf{y}}^T] = \mathbf{H} \mathbb{E}[\mathbf{w} \mathbf{w}^T] \mathbf{H}^T = \sigma^2 \mathbf{H} \mathbf{H}^T = \mathbf{K}$$

Gaussian processes

- A Gaussian process is defined as probability distribution over functions $g(\mathbf{x})$
- such that the set of values of $g(\mathbf{x})$ evaluated at a set $\{\mathbf{x}^1, \dots, \mathbf{x}^n\}$ is jointly a Gaussian distribution.
- Gaussian distribution is fully determined by its first two moments, **mean** and **covariance**
- The mean vector must be assumed to be zero 0, such that the specification is mostly determined by the **covariance**
- Covariance is determined by the kernel function and thus the pairwise similarity of data points in some space

Connection of neural networks and Gaussian processes

- Discovered by Neal (1994) in his PhD thesis as pessimistic result
- Two-layer network with weights $v_{jd} \sim \mathcal{N}(0, \sigma_w^2)$ and $w_{kj} \sim \mathcal{N}(0, \sigma_w^2)$.
- Pre-activations have identical means and second moments

$$\mathbb{E}[\sum_{d=1}^D x_d v_{jd}] = 0 \quad \mathbb{E}[s_j^2(\mathbf{x})] = \mathbb{E}[\sum_{d=1}^D x_d^2 v_{jd}^2] = \mathbb{E}[s_l^2(\mathbf{x})]$$

- Activations in hidden layer have some mean m_h and variance V_h .
- Distributions of outputs:

$$\mathbb{E}[o_k(\mathbf{x})] = \mathbb{E}[\sum_{j=1}^J w_{kj} h_j(\mathbf{x})] = J \mathbb{E}[w_{kj}] \mathbb{E}[h_1(\mathbf{x})] = 0.$$

$$\mathbb{E}[(o_k(\mathbf{x}))^2] = \mathbb{E}[\sum_{j=1}^J (w_{kj} h_j(\mathbf{x}))^2] = J \mathbb{E}[w_{kj}^2] \mathbb{E}[(h_1(\mathbf{x}))^2] = J \sigma_w^2 V_h.$$

Connection of neural networks and Gaussian processes

- We can use the *Central Limit Theorem* to show that outputs are Gaussian distributed (similar to our derivation of SNNs)
- The joint distribution of two data points then is

$$\mathbb{E}[o_k(\mathbf{x}_n)o_k(\mathbf{x}_m)] = \sum_{j=1}^J \mathbb{E}[w_{jk}h_j(\mathbf{x}_n)w_{jk}h_j(\mathbf{x}_m)] = \sigma_w^2 V_h K(\mathbf{x}_n, \mathbf{x}_m)$$

- Gaussian for multiple data points, thus a Gaussian process whose covariance depends on kernel between data points
- Covariance between two different output units is zero
- Connection is often used for Deep Learning theory
 - Also for *Neural Tangent Kernels*